



# Tutorial: Partitioning, Load Balancing and the Zoltan Toolkit

**Erik Boman and Karen Devine  
Discrete Algorithms and Math Dept.  
Sandia National Laboratories, NM**

**CSCAPES Institute**

**SciDAC Tutorial, MIT, June 2007**



Sandia is a multiprogram laboratory operated by Sandia Corporation, a Lockheed Martin Company,  
for the United States Department of Energy's National Nuclear Security Administration  
under contract DE-AC04-94AL85000.





# Outline

---

## Part 1:

- Partitioning and load balancing
  - “Owner computes” approach
- Static vs. dynamic partitioning
- Models and algorithms
  - Geometric (RCB, SFC)
  - Graph & hypergraph

## Part 2:

- Zoltan
  - Capabilities
  - How to get it, configure, build
  - How to use Zoltan with your application

# Parallel Computing in CS&E

- **Parallel Computing Challenge**
  - **Scientific simulations critical to modern science.**
    - Models grow in size, higher fidelity/resolution.
    - Simulations must be done on parallel computers.
  - **Clusters with 64-256 nodes are widely available.**
  - **High-performance computers have 100,000+ processors.**
    - How can we use such machines efficiently?



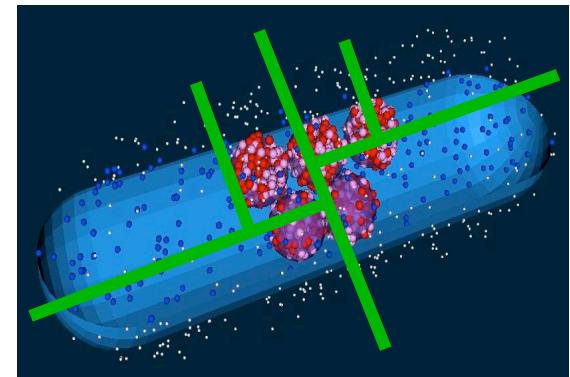
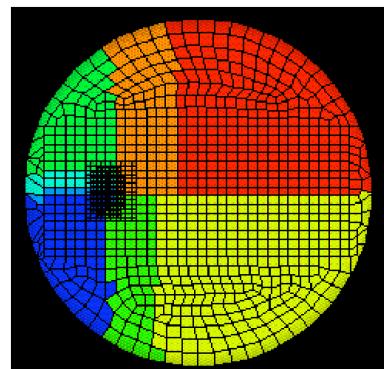
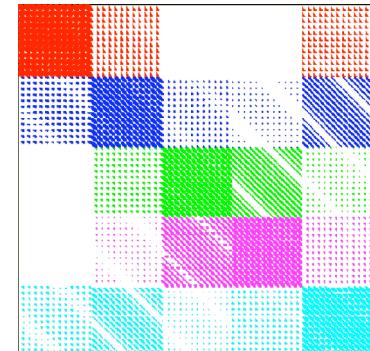
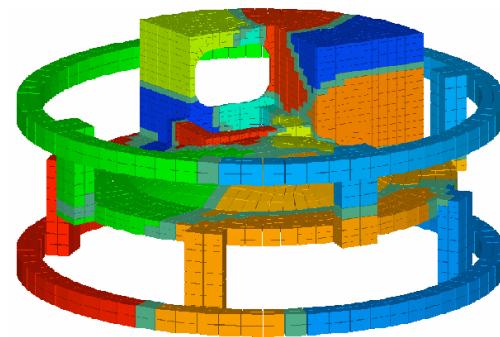
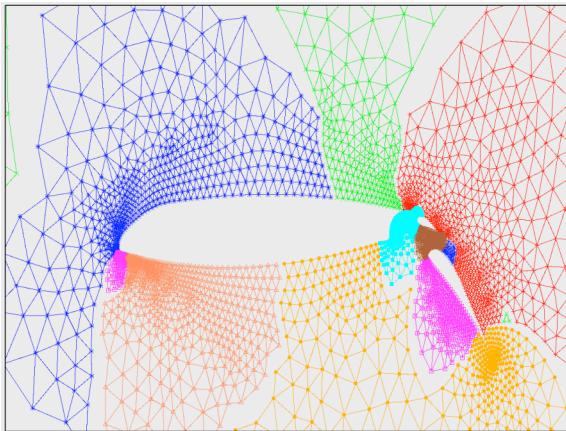
# Parallel Computing Approaches

---

- We focus on distributed memory systems.
  - Two common approaches:
- Master–slave
  - A “master” processor is a global synchronization point, hands out work to the slaves.
- Data decomposition + “Owner computes”:
  - The data is distributed among the processors.
  - The owner performs all computation on its data.
  - Data distribution defines work assignment.
  - Data dependencies among data items owned by different processors incur communication.

# Partitioning and Load Balancing

- Assignment of application data to processors for parallel computation.
- Applied to grid points, elements, matrix rows, particles, ....

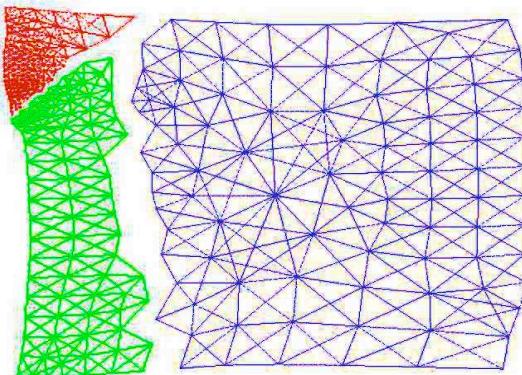




# Partitioning Goals

---

- **Minimize total execution time by...**
  - **Minimizing processor idle time.**
    - Load balance data and work.
  - **Keeping inter-processor communication low.**
    - Reduce total volume, max volume.
    - Reduce number of messages.



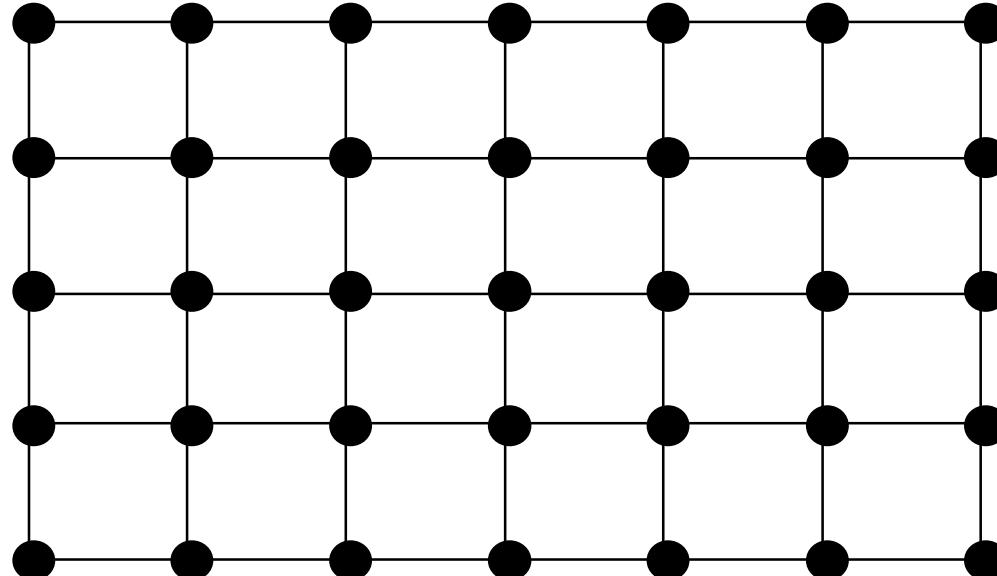
*Partition of an unstructured finite element mesh for three processors*



# “Simple” Example (1)

---

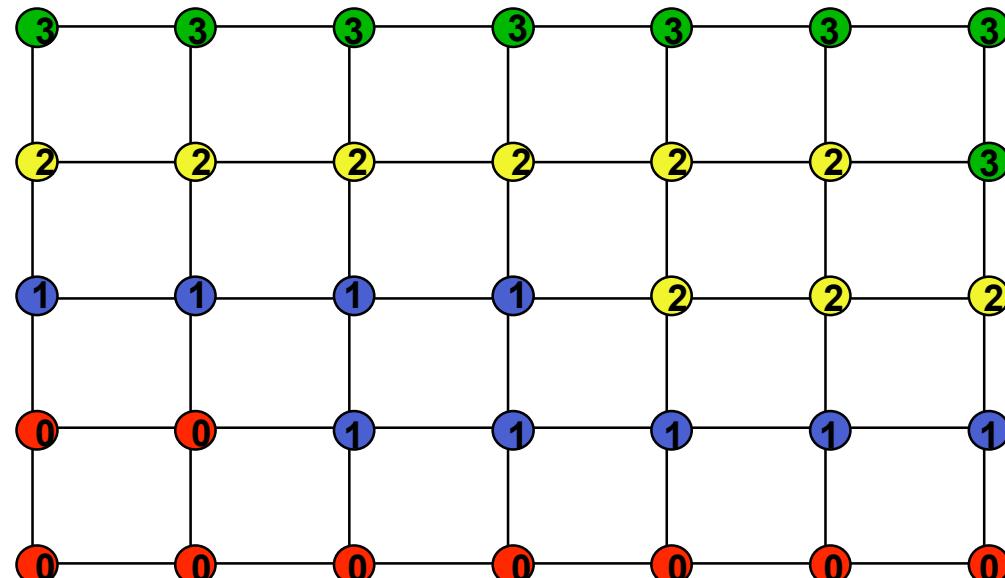
- Finite difference method.
  - Assign equal numbers of grid points to processors.
  - Keep amount of data communicated small.



7x5 grid  
5-point stencil  
4 processors

# “Simple” Example (2)

- Finite difference method.
  - Assign equal numbers of grid points to processors.
  - Keep amount of data communicated small.



Max Data Comm: 14

Total Volume: 42

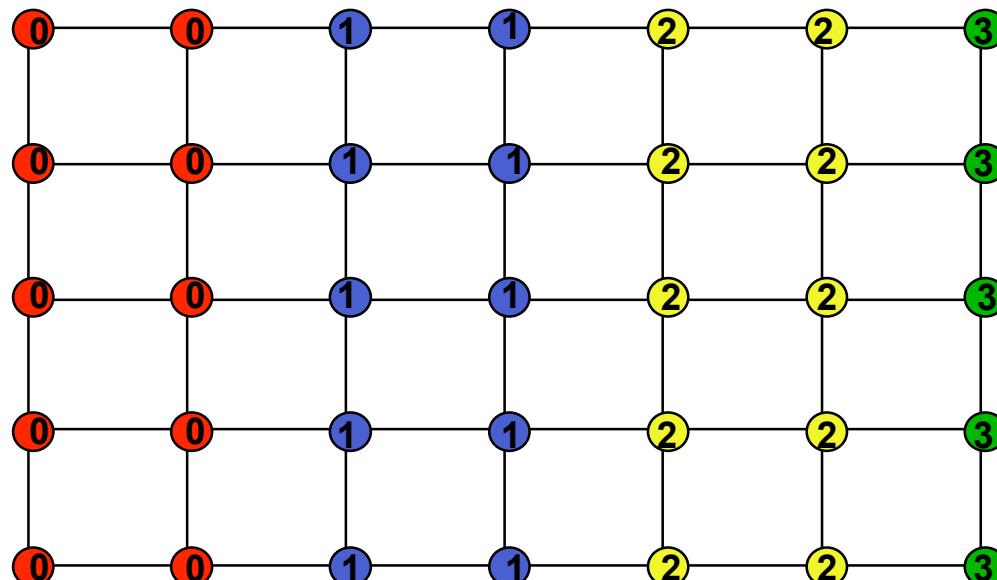
Max Nbor Proc: 2

Max Imbalance: 3%

*First 35/4 points to processor 0;  
next 35/4 points to processor 1; etc.*

# “Simple” Example (3)

- Finite difference method.
  - Assign equal numbers of grid points to processors.
  - Keep amount of data communicated small.



Max Data Comm: 10

Total Volume: 30

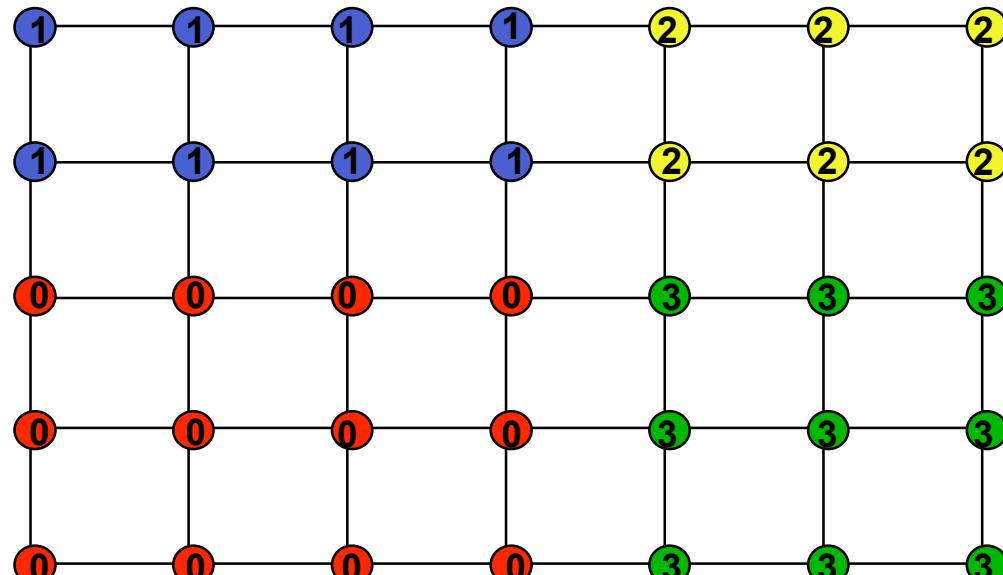
Max Nbor Proc: 2

Max Imbalance: 14%

*One-dimensional striped partition*

# “Simple” Example (4)

- Finite difference method.
  - Assign equal numbers of grid points to processors.
  - Keep amount of data communicated small.



Max Data Comm: 7

Total Volume: 26

Max Nbor Proc: 2

Max Imbalance: 37%

*Two-dimensional  
structured grid partition*



# Static Partitioning

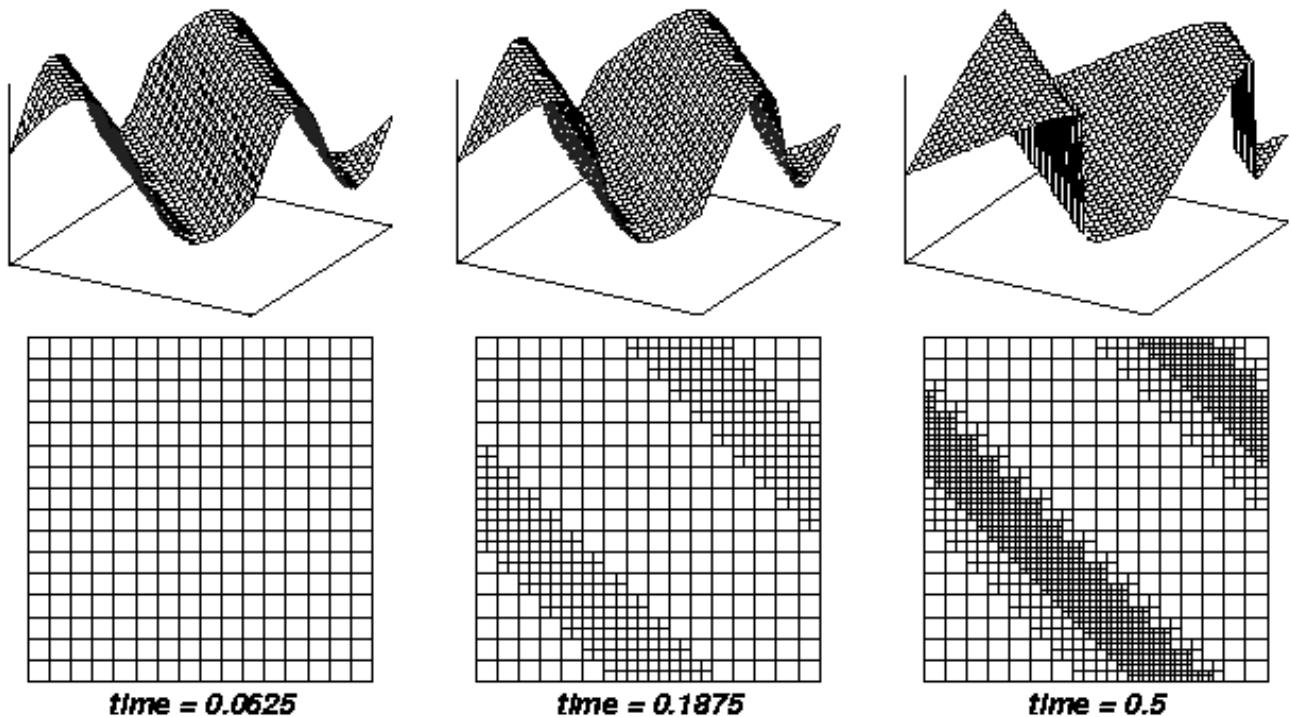


- **Static partitioning in an application:**
  - Data partition is computed.
  - Data are distributed according to partition map.
  - Application computes.
- **Ideal partition:**
  - Processor idle time is minimized.
  - Inter-processor communication costs are kept low.

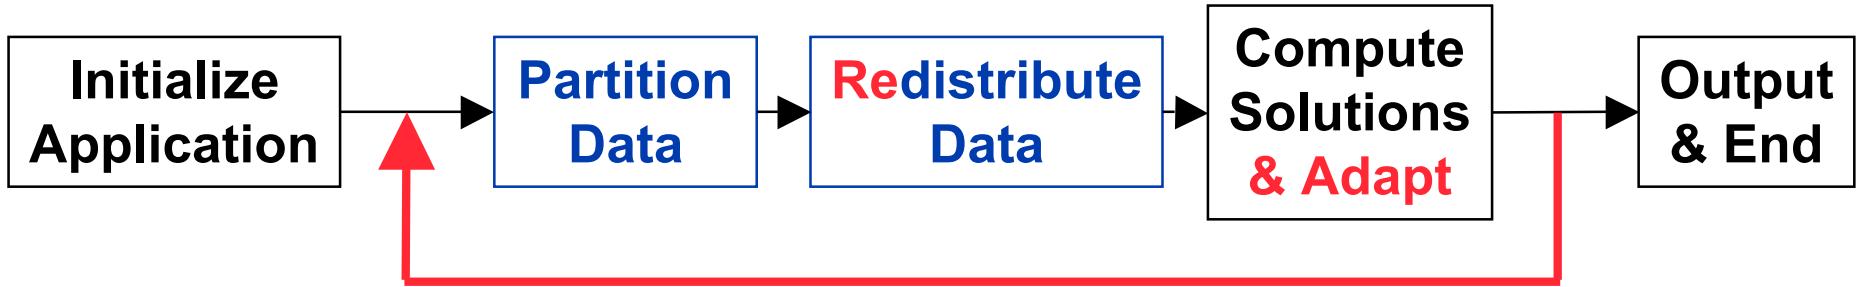
# Dynamic Applications

- Characteristics:
  - Work per processor is unpredictable or changes during a computation; and/or
  - Locality of objects changes during computations.
  - Dynamic redistribution of work is needed during computation.

- Example:  
adaptive  
mesh  
refinement  
(AMR)  
methods



# Dynamic Repartitioning (a.k.a. Dynamic Load Balancing)



- Dynamic repartitioning (load balancing) in an application:
  - Data partition is computed.
  - Data are distributed according to partition map.
  - Application computes and, perhaps, adapts.
  - Process repeats until the application is done.
- Ideal partition:
  - Processor idle time is minimized.
  - Inter-processor communication costs are kept low.
  - Cost to redistribute data is also kept low.



# Static vs. Dynamic:

## Usage and Implementation

---

Slide 14



Sandia  
National  
Laboratories

- **Static:**

- Pre-processor to application.
- Can be implemented serially.
- May be slow, expensive.
- File-based interface acceptable.
- No consideration of existing decomposition required.

- **Dynamic:**

- Must run side-by-side with application.
- Must be implemented in parallel.
- Must be fast, scalable.
- Library application interface required.
- Should be easy to use.
- Incremental algorithms preferred.
  - Small changes in input result small changes in partitions.
  - Explicit or implicit incrementality acceptable.

# Two Types of Models/Algorithms

---

- **Geometric**
  - Computations are tied to a geometric domain.
  - Coordinates for data items are available.
  - Geometric locality is loosely correlated to data dependencies.
- **Combinatorial (topological)**
  - No geometry .
  - Connectivity among data items is known.
    - Represent as graph or hypergraph.



# Recursive Coordinate Geometric Bisection (RCB)

Slide 16

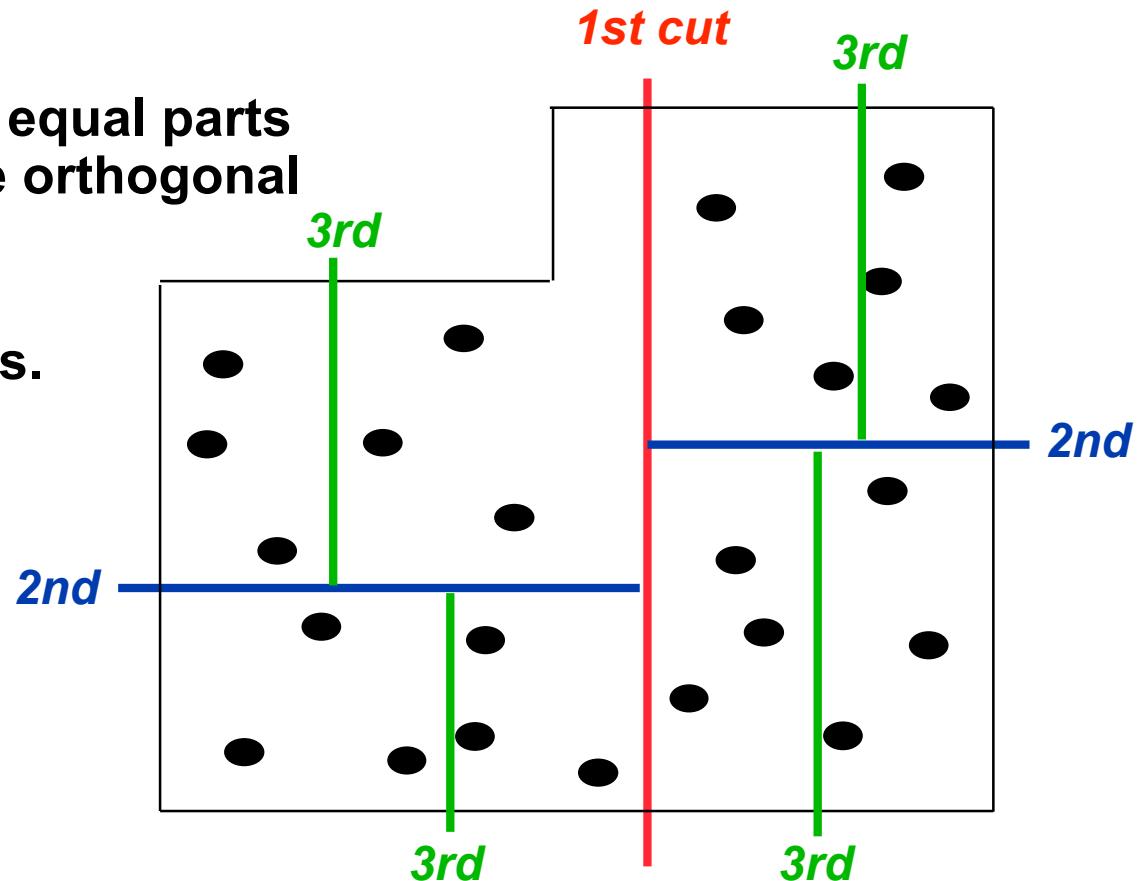


- Developed by Berger & Bokhari (1987) for AMR.

- Independently discovered by others.

- Idea:

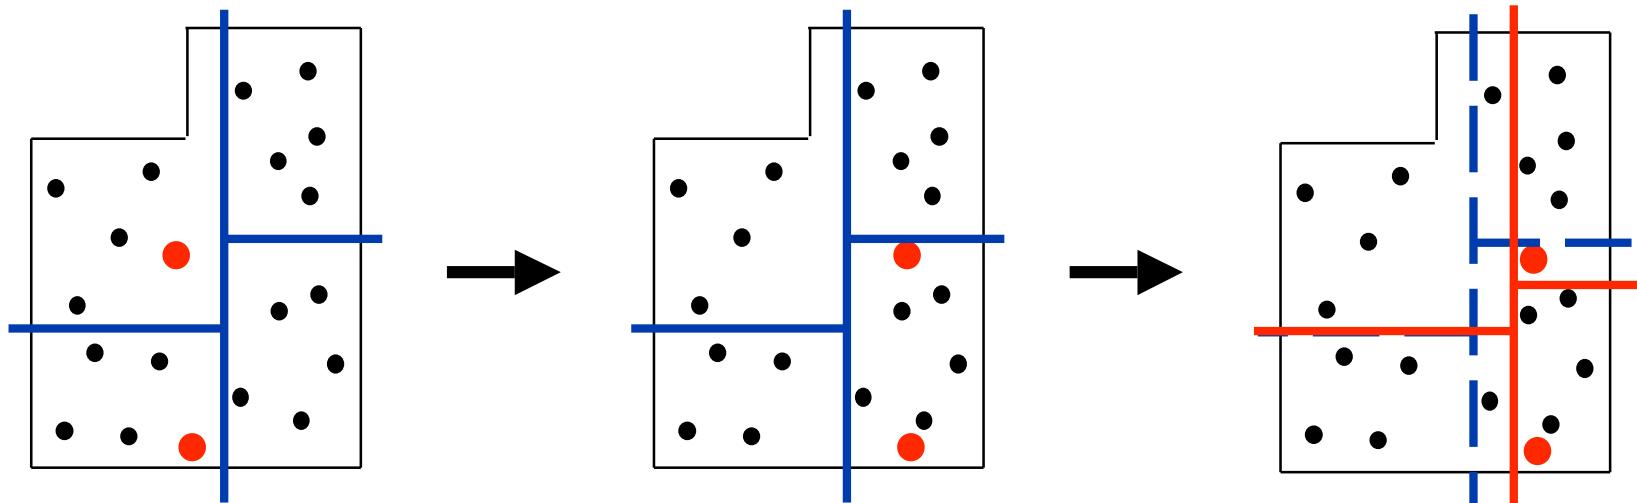
- Divide work into two equal parts using a cutting plane orthogonal to a coordinate axis.
  - Recursively cut the resulting subdomains.





# RCB Repartitioning

- **Implicitly incremental.**
- **Small changes in data results in small movement of cuts.**





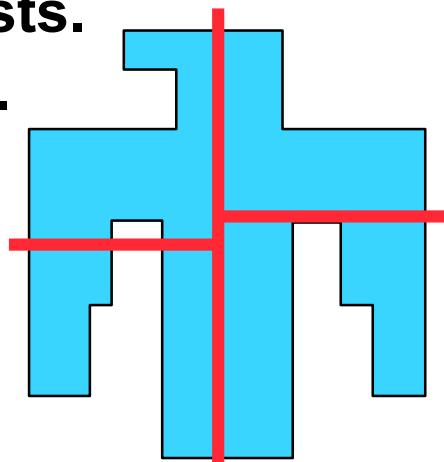
# RCB Advantages and Disadvantages

---

Slide 18

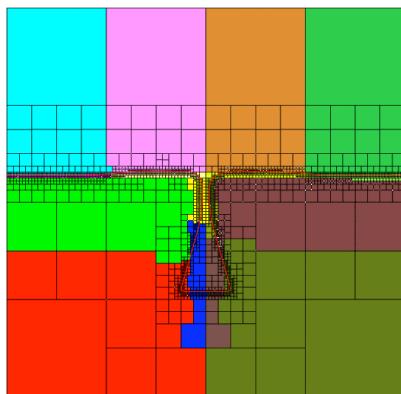


- **Advantages:**
  - Conceptually simple; fast and inexpensive.
  - Regular subdomains.
    - Can be used for structured or unstructured applications.
    - All processors can inexpensively know entire decomposition.
  - Effective when connectivity info is not available.
- **Disadvantages:**
  - No explicit control of communication costs.
  - Can generate disconnected subdomains.
  - Mediocre partition quality.
  - Geometric coordinates needed.

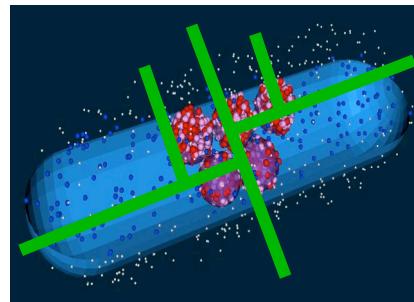




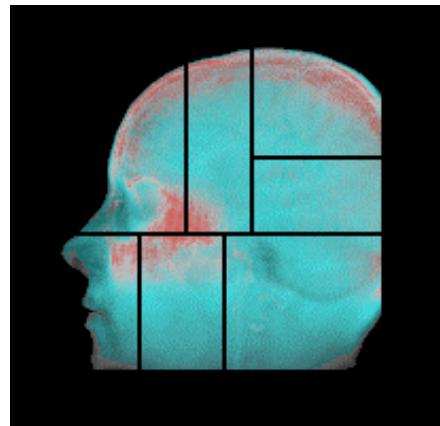
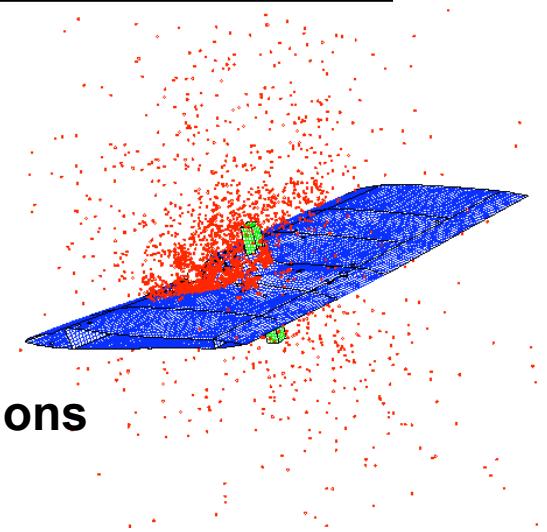
# Applications of RCB



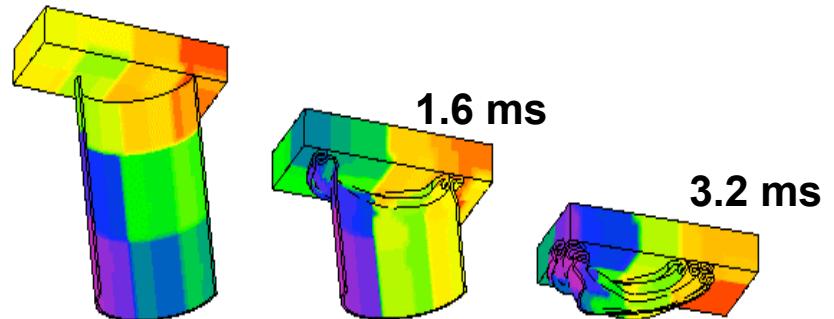
Adaptive Mesh Refinement



Particle Simulations



Parallel Volume Rendering



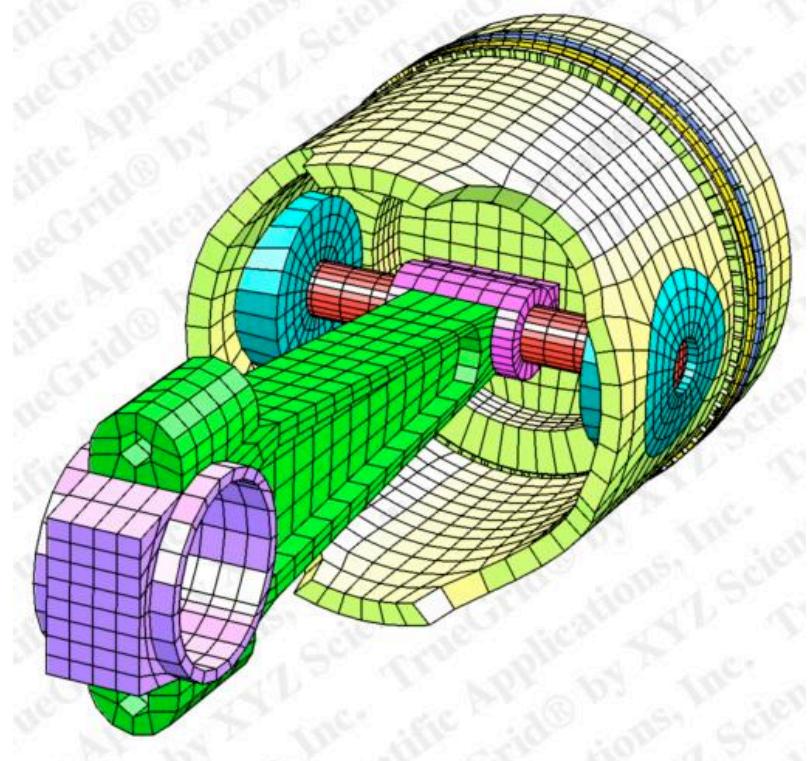
Crash Simulations  
and Contact Detection



# Variations on RCB : RIB

---

- **Recursive Inertial Bisection**
  - Simon, Taylor, et al., 1991
  - Cutting planes orthogonal to principle axes of geometry.
  - Not incremental.



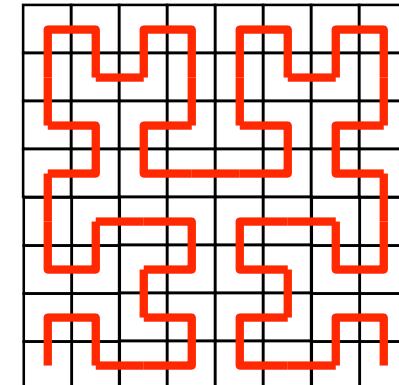
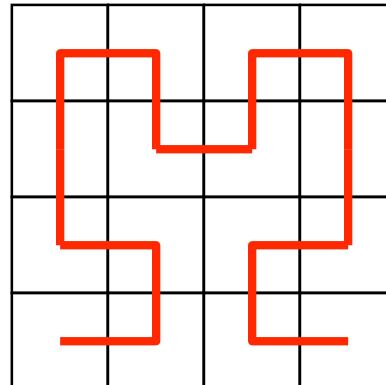
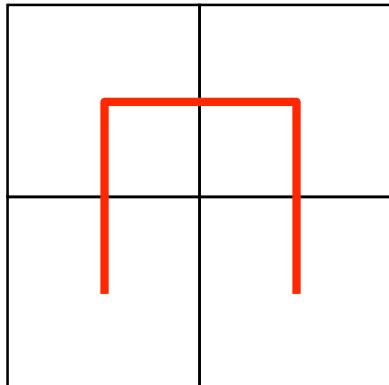


# Space-Filling Curve Partitioning (SFC)

Slide 21



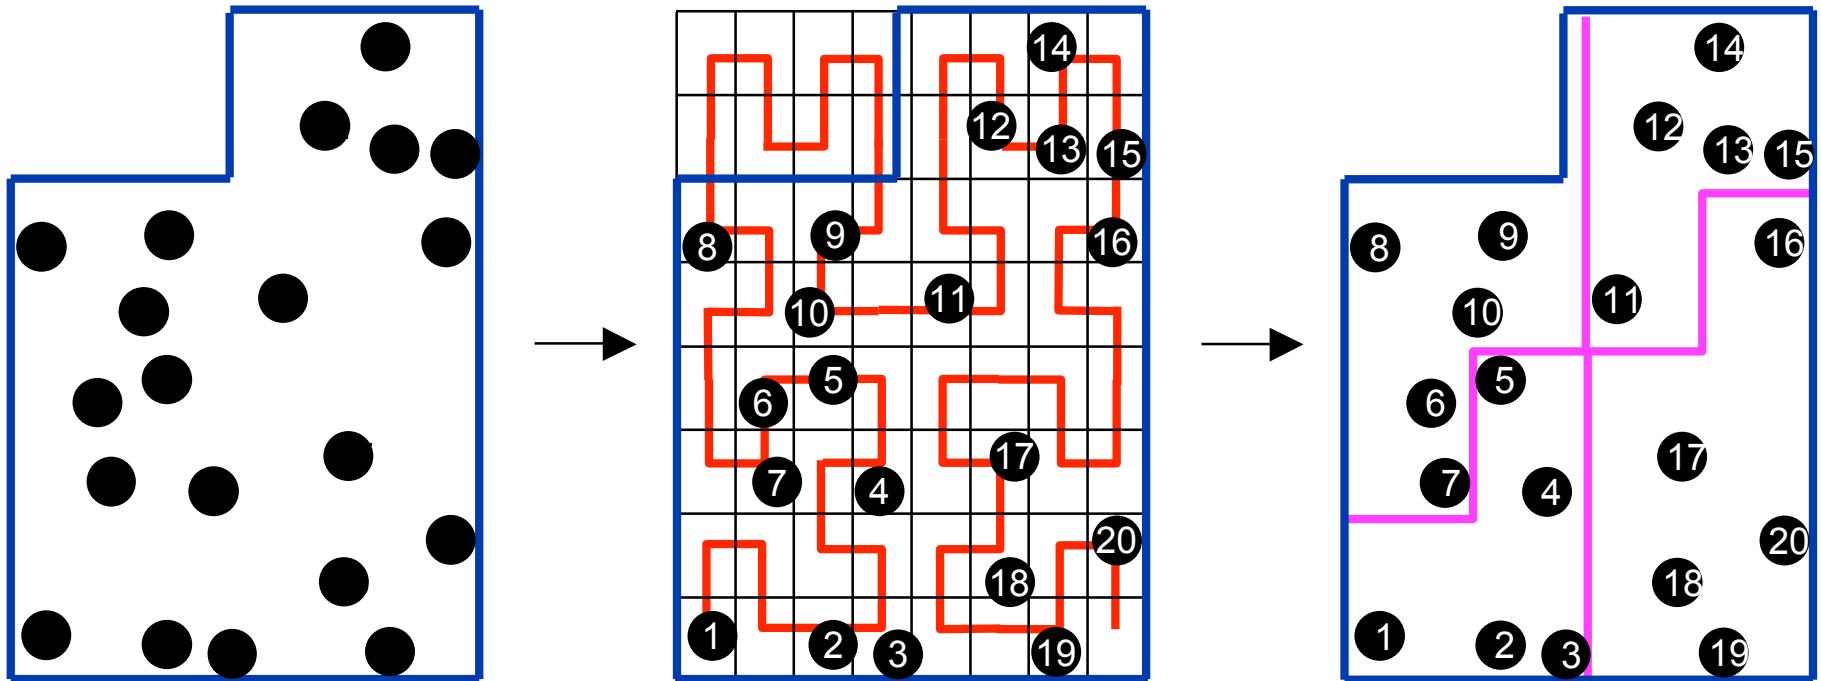
- Developed by Peano, 1890.
- Space-Filling Curve:
  - Mapping between  $R^3$  to  $R^1$  that completely fills a domain.
  - Applied recursively to obtain desired granularity.
- Used for partitioning by ...
  - Warren and Salmon, 1993, gravitational simulations.
  - Pilkington and Baden, 1994, smoothed particle hydrodynamics.
  - Patra and Oden, 1995, adaptive mesh refinement.





# SFC Algorithm

- Run space-filling curve through domain.
- Order objects according to position on curve.
- Perform 1-D partition of curve.

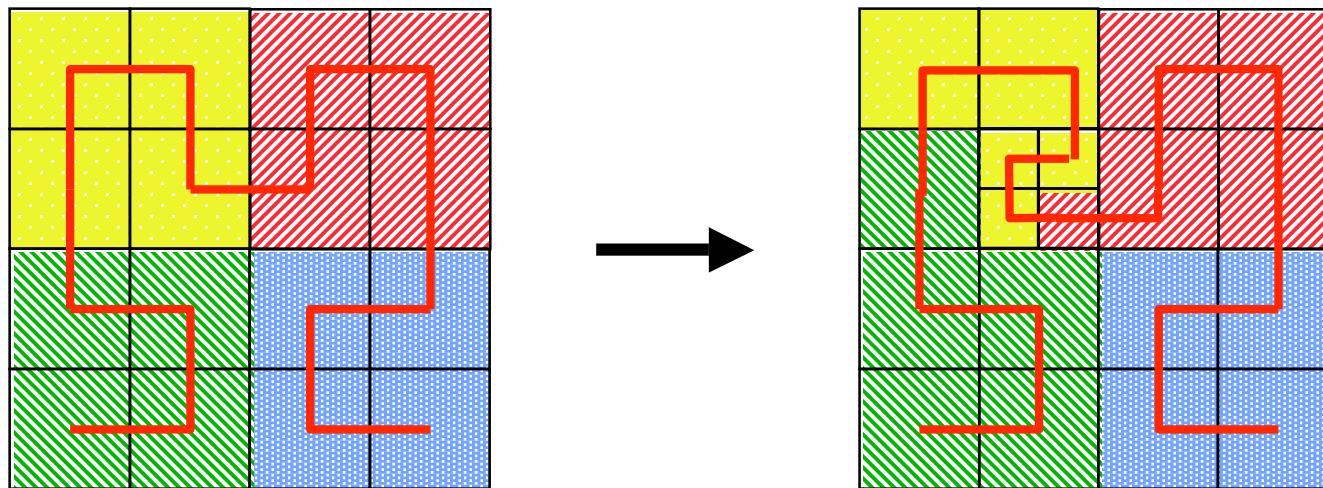




# SFC Repartitioning

---

- Implicitly incremental.
- Small changes in data results in small movement of cuts in linear ordering.





# SFC Advantages and Disadvantages

---

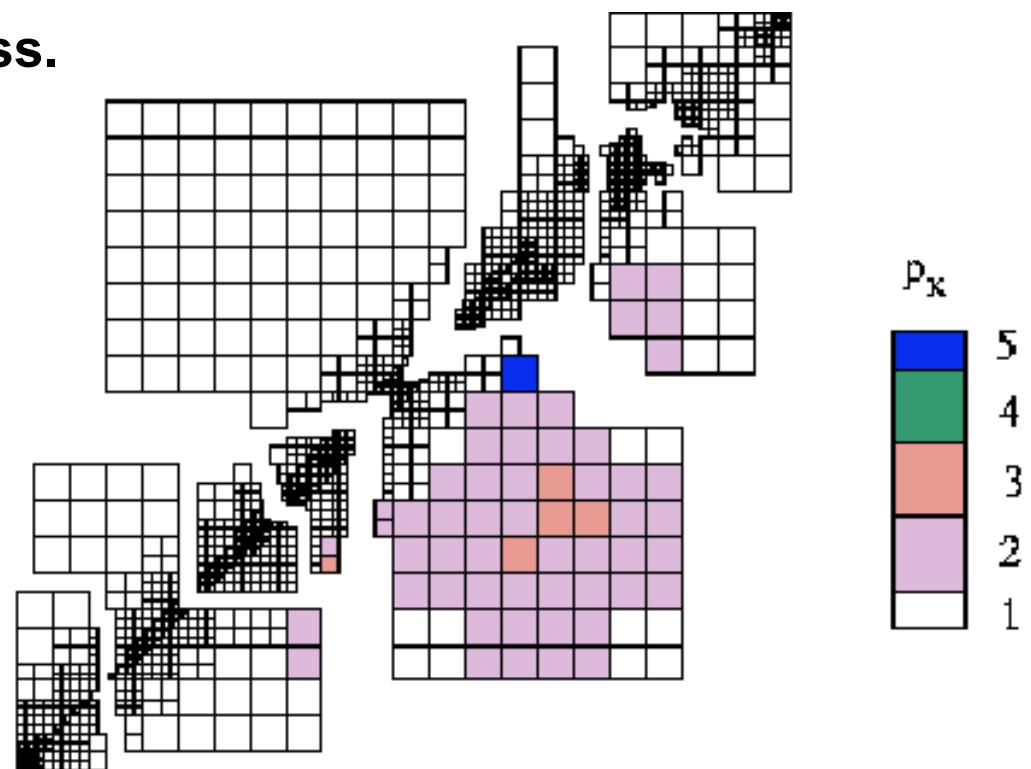
Slide 24



- **Advantages:**
  - Simple, fast, inexpensive.
  - Maintains geometric locality of objects in processors.
  - Linear ordering of objects may improve cache performance.
- **Disadvantages:**
  - No explicit control of communication costs.
  - Can generate disconnected subdomains.
  - Often lower quality partitions than RCB.
  - Geometric coordinates needed.

# Applications using SFC

- Adaptive hp-refinement finite element methods.
  - Assigns physically close elements to same processor.
  - Inexpensive; incremental; fast.
  - Linear ordering can be used to order elements for efficient memory access.



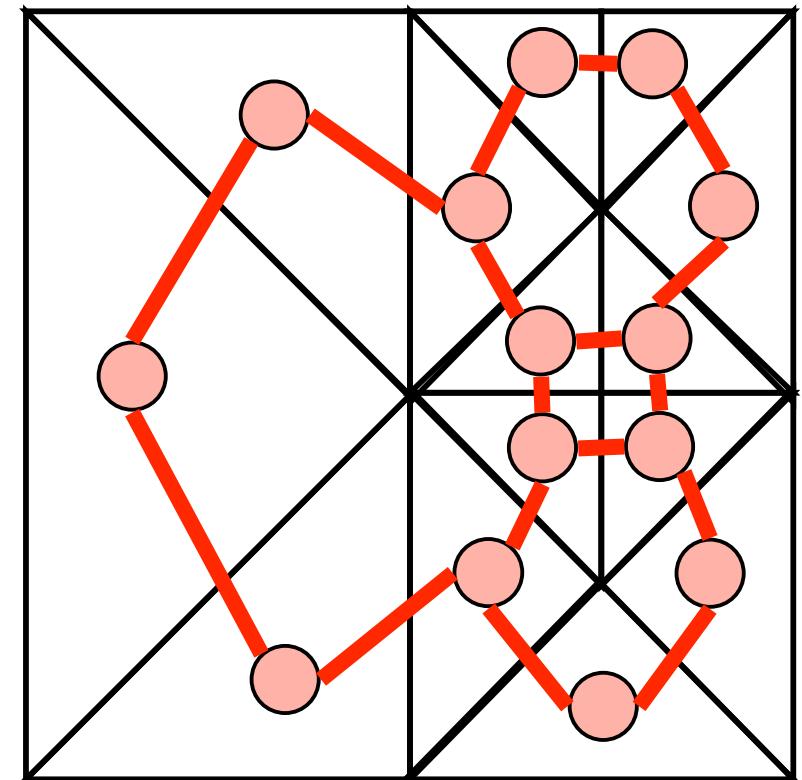
*hp-refinement mesh; 8 processors.  
Patra, et al. (SUNY-Buffalo)*



# Graph Partitioning

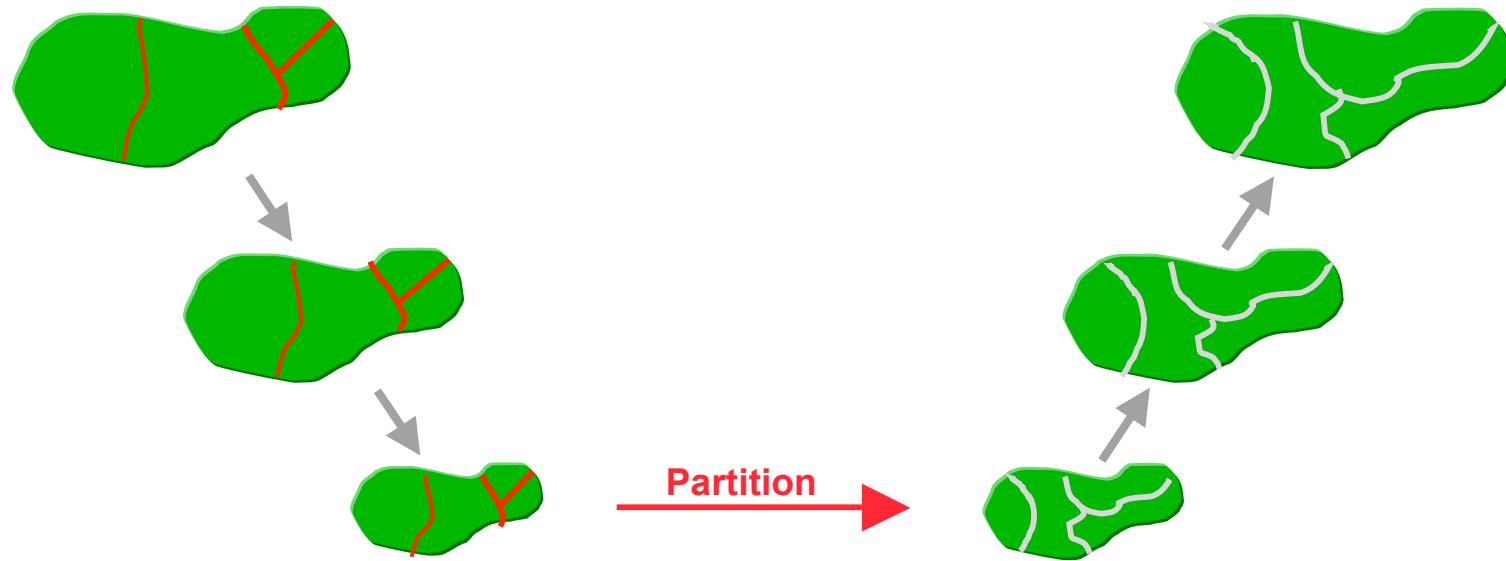
---

- Represent problem as a weighted graph.
  - Vertices = objects to be partitioned.
  - Edges = communication between objects.
  - Weights = work load or amount of communication.
- Partition graph so that ...
  - Partitions have equal vertex weight.
  - Weight of edges cut by subdomain boundaries is small.



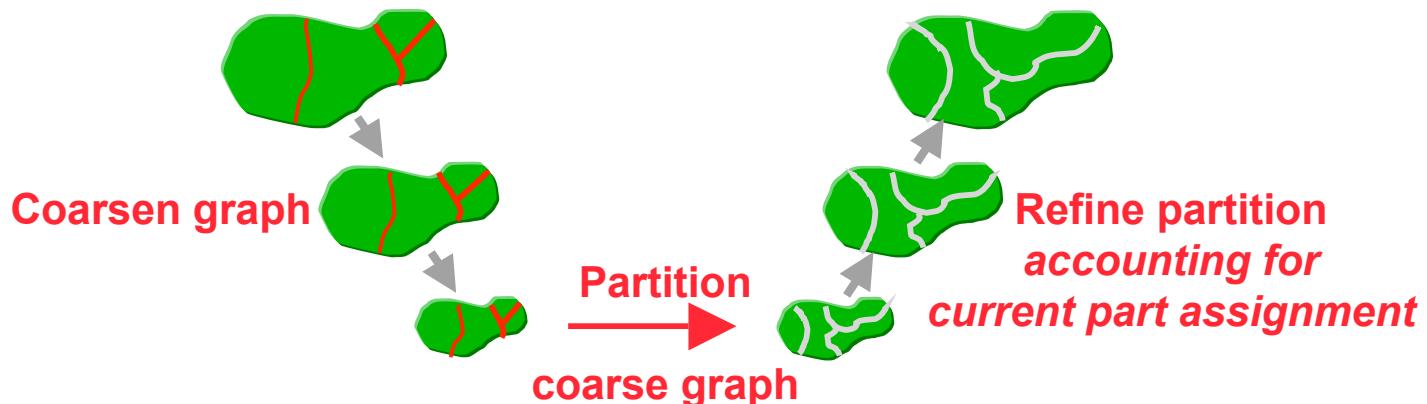
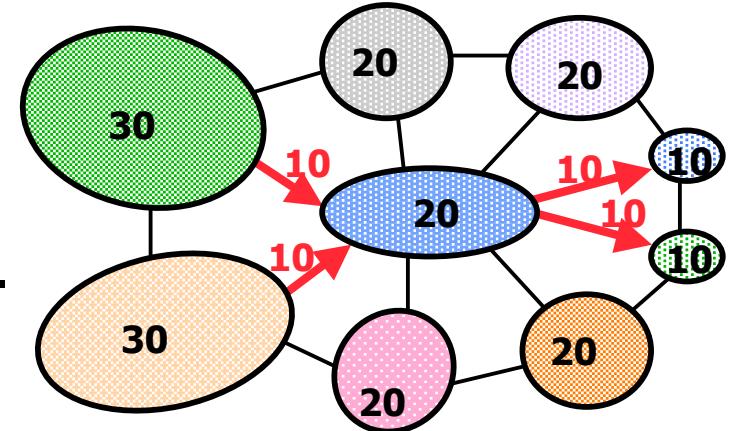
# Multi-Level Graph Partitioning

- Bui & Jones (1993); Hendrickson & Leland (1993); Karypis and Kumar (1995)
- Construct smaller approximations to graph.
- Perform graph partitioning on coarse graph.
- Propagate partition back, refining as needed.



# Graph Repartitioning

- Diffusive strategies (Cybenko, Hu, Blake, Walshaw, Schloegel, et al.)
  - Shift work from highly loaded processors to less loaded neighbors.
  - Local communication keeps data redistribution costs low.
- Multilevel partitioners that account for data redistribution costs in refining partitions (Schloegel, Karypis)
  - Parameter weights application communication vs. redistribution communication.





# Graph Partitioning

## Advantages and Disadvantages

Slide 29



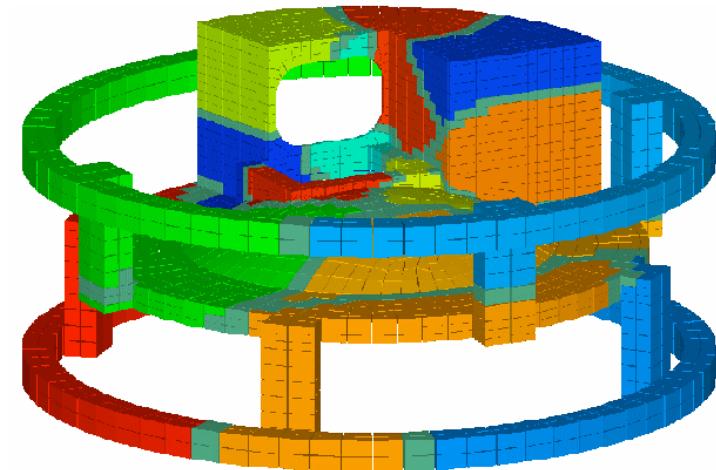
Sandia  
National  
Laboratories

- **Advantages:**

- High quality partitions for many applications.
- Explicit control of communication costs.
- Widely used for static partitioning (Chaco, METIS, Jostle, Party, Scotch)

- **Disadvantages:**

- More expensive than geometric approaches.
- Not incremental.





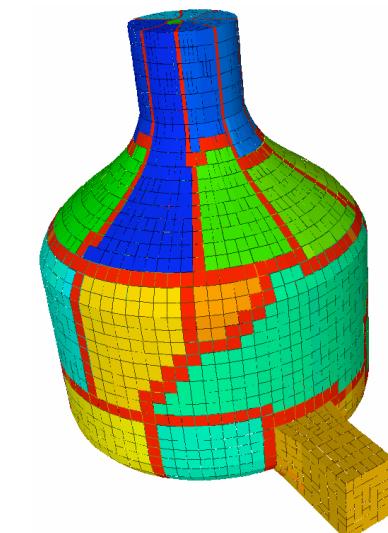
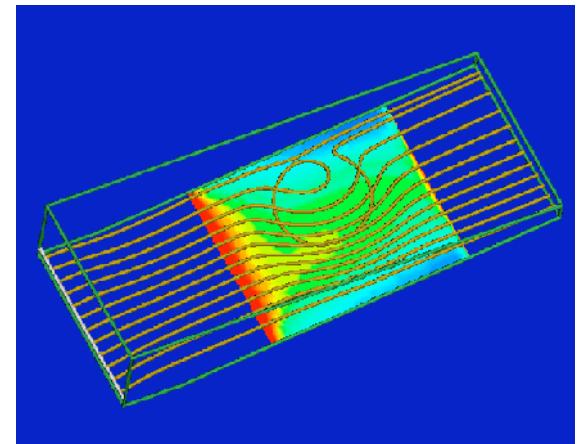
# Applications using Graph Partitioning

---

Slide 30



- Finite element analysis
- Multiphysics simulations
  - Difficult to estimate work in advance.
  - Rebalance infrequently; want high quality.
- Linear solvers and preconditioners
  - Square, structurally symmetric.
  - Decomposition of mesh induces good decomposition for solver.



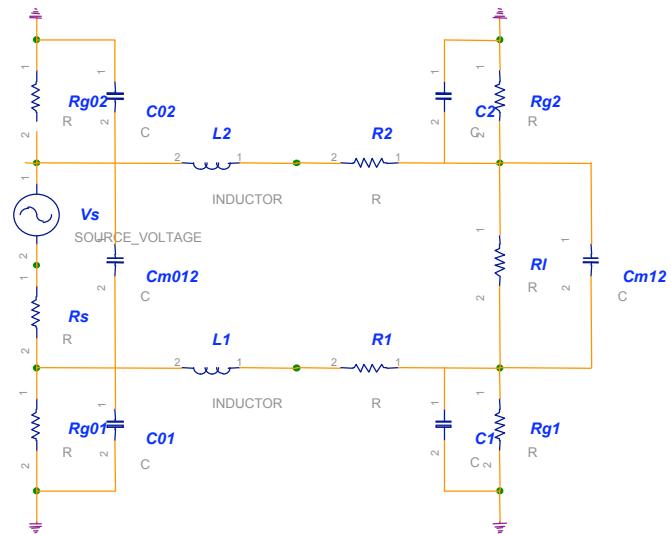


# Applications using Graph Partitioning

Slide 31

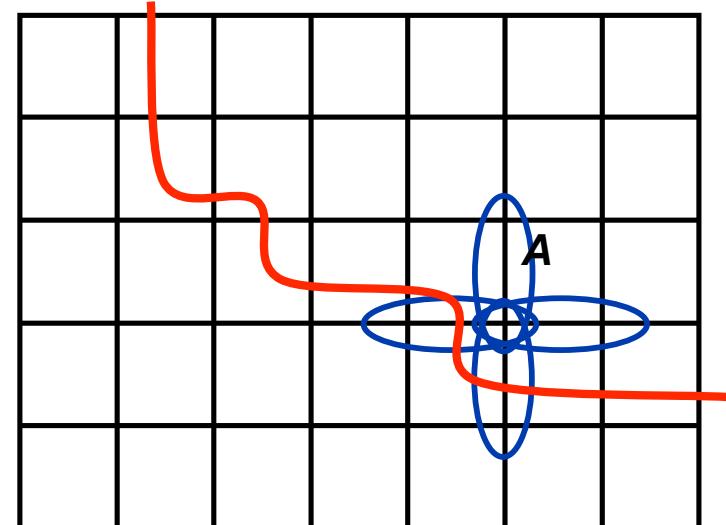


- XYCE (S. Hutchinson, R. Hoekstra, E. Keiter, SNL)
  - Massively parallel analog circuit simulator.
- Load balancing in XYCE.
  - Balance linear solve phase.
  - Equal number of rows while minimizing cut edges.
  - Partition matrix solver separately from matrix fill.
  - Trilinos solver library (Heroux, et al.) uses Zoltan to partition matrix.
- Matrix structure more complex than mesh-based applications.
  - Is there a better partitioning model?



# Flaws in the Graph Model

- Graph model and partitioning approach has been successful in scientific computing, BUT...
- Graph models assume # edge cuts = communication volume.
- In reality...
  - Edge cuts are not equal to communication volume.





# Graph Models: Applicability

---

- Graph models assume symmetric square problem.
  - Symmetric == undirected graph.
  - Square == inputs and outputs of operation are same size.
- Non-symmetric systems.
  - Require directed or bipartite graph.
- Rectangular systems.
  - Require decompositions for differently sized inputs and outputs.

$$\begin{matrix} \textcolor{magenta}{y} \\ \textcolor{magenta}{=} \\ \textcolor{magenta}{A} \\ \textcolor{magenta}{x} \end{matrix}$$

A diagram illustrating a rectangular system. On the left is a vertical vector  $y$  represented by a column of colored squares (magenta). In the center is a matrix  $A$  represented by a grid of colored squares (white and red). On the right is a vertical vector  $x$  represented by a column of colored squares (green).

$$\begin{matrix} \textcolor{magenta}{y} \\ \textcolor{magenta}{=} \\ \textcolor{magenta}{A} \\ \textcolor{magenta}{x} \end{matrix}$$

A diagram illustrating a rectangular system. On the left is a vertical vector  $y$  represented by a column of colored squares (magenta). In the center is a matrix  $A$  represented by a grid of colored squares (white, red, and pink). On the right is a vertical vector  $x$  represented by a column of colored squares (green and pink).

# Is the Graph Model “good enough”?

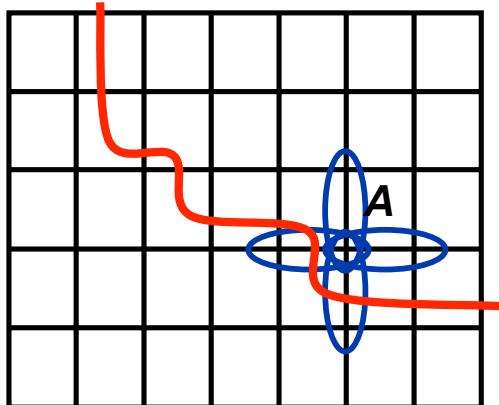
---

- Mesh-based applications: Yes, maybe.
  - Graph partitioning works well in practice.
  - Geometric structure of meshes ensures...
    - Small separators and good partitions.
    - Low vertex degrees give small error in graph model.
- Irregular non-mesh applications: No.
  - Graph model is poor or does not apply.
  - Ex: circuit simulation, discrete optimization, data mining.
  - Nonsymmetric and rectangular matrices.

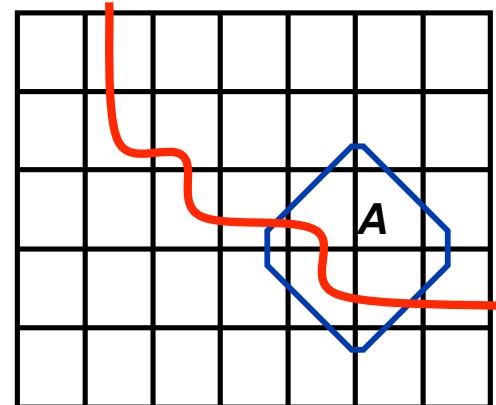


# Hypergraph Partitioning

- **Hypergraph model** (Aykanat & Catalyurek)
  - Vertices represent computations.
  - Hyperedges connect all objects which produce/use datum.
    - Hyperedges connect one or more vertices (cf. graph edge always two)
  - Greater expressiveness than graph partitioners.
    - Non-symmetric data dependencies.
    - Rectangular matrices.
  - Cut hyperedges == communication volume!



*Graph model only approximates communication volume.*

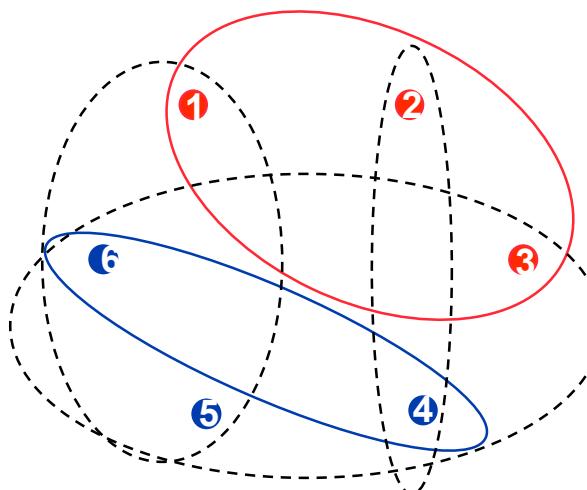


*Hypergraph model accurately measures communication volume.*

# Matrices and Hypergraphs

- View matrix as hypergraph (Çatalyürek & Aykanat)
  - Vertices == columns
  - Edges == rows
- Partition vertices (columns in matrix)
- Communication volume associated with edge e:
 
$$CV_e = (\# \text{ processors in edge } e) - 1$$
- Total communication volume =

$$\sum_e CV_e$$



$$\begin{pmatrix} \mathbf{y} \\ \mathbf{y} \\ \mathbf{y} \\ \mathbf{y} \\ \mathbf{y} \end{pmatrix} = \begin{pmatrix} * & * & * \\ & * & * \\ * & & * \\ & * & * \\ * & * & * \\ & * & * \end{pmatrix} \begin{pmatrix} \mathbf{x} \\ \mathbf{x} \\ \mathbf{x} \\ \mathbf{x} \\ \mathbf{x} \end{pmatrix}$$

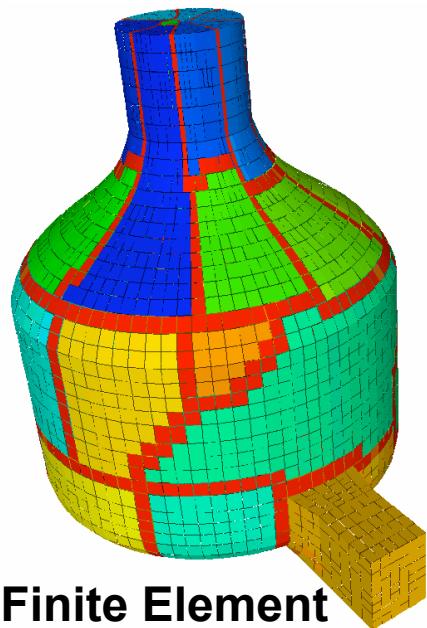


# Hypergraph Repartitioning

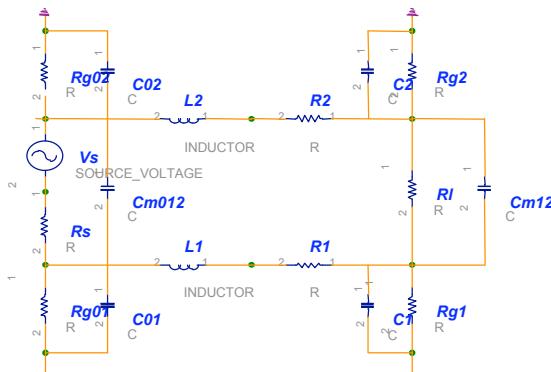
---

- Augment hypergraph with data redistribution costs.
  - Account for data's current processor assignments.
  - Weight hyperedges by data redistribution size or frequency of use.
- Hypergraph partitioning then attempts to minimize *total communication volume*:  
$$\begin{aligned} &\text{Data redistribution volume} \\ &+ \text{Application communication volume} \\ &\text{Total communication volume} \end{aligned}$$
- Trade-off between application volume and redistribution cost controlled by single knob (user parameter).
  - PHG\_REPART\_MULTIPLIER should be (roughly) number of application communications between repartitions.
- Can re-use existing algorithms and software.
  - This approach also works for graphs.

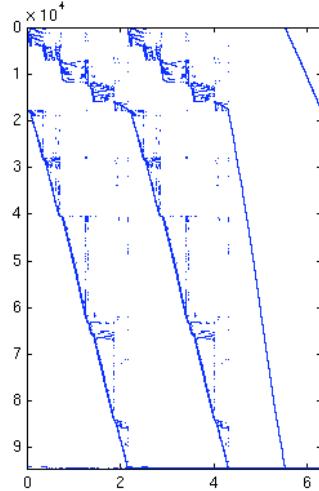
# Hypergraph Applications



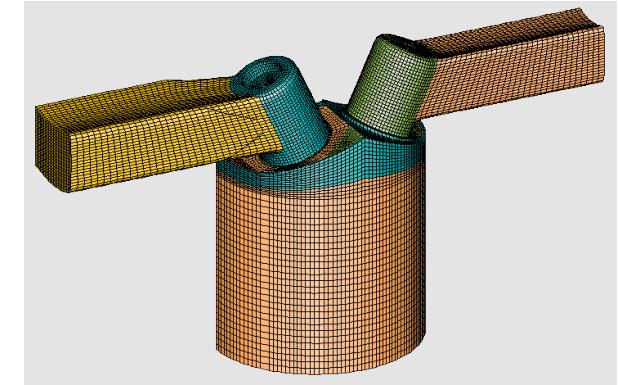
**Finite Element Analysis**



**Circuit Simulations**



**Linear programming  
for sensor placement**

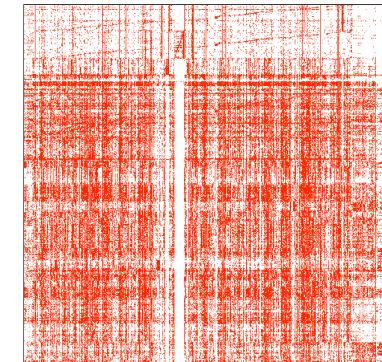


**Multiphysics and  
multiphase simulations**

$$\begin{matrix} & \text{A} & \\ \text{x} & = & \text{b} \end{matrix}$$

The matrix A is a 10x10 grid with red blocks forming a diagonal pattern. The vector x is a 10x1 column vector with alternating red and white entries. The vector b is a 10x1 column vector with alternating red and white entries.

**Linear solvers & preconditioners  
(no restrictions on matrix structure)**



**Data Mining**



# Hypergraph Partitioning: Advantages and Disadvantages

---

Slide 39



Sandia  
National  
Laboratories

- **Advantages:**
  - Communication volume reduced 30-38% on average over graph partitioning (Catalyurek & Aykanat).
    - 5-15% reduction for mesh-based applications.
  - More accurate communication model than graph partitioning.
    - Better representation of highly connected and/or non-homogeneous systems.
  - Greater applicability than graph model.
    - Can represent rectangular systems and non-symmetric dependencies.
- **Disadvantages:**
  - More expensive than graph partitioning.



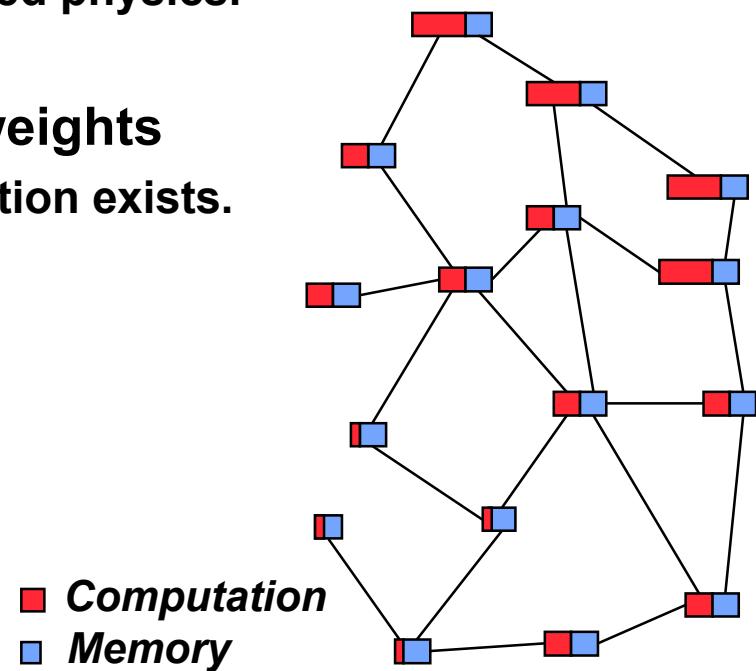
# Using Weights

---

- Some data items may have more work than others.
- Solution: Specify work (load) using weights.
  - By default, all data items have unit weights.
  - Objective is to balance sum of weights.
- Geometric methods:
  - Add a weight for each point.
- Graph/hypergraph methods:
  - One weight per vertex.
  - Can also weight edges with communication size.

# Multi-criteria Load-balancing

- Multiple constraints or objectives
  - Compute a single partition that is good with respect to multiple factors.
    - Balance both computation and memory.
    - Balance meshes in loosely coupled physics.
    - Balance multi-phase simulations.
  - Extend algorithms to multiple weights
    - Difficult. No guarantee good solution exists.



# Heterogeneous Architectures

---

- Clusters may have different types of processors.
- Assign “capacity” weights to processors.
  - Compute power (speed)
  - Memory
- Partitioner should balance with respect to processor capacity.



## Example & Recap

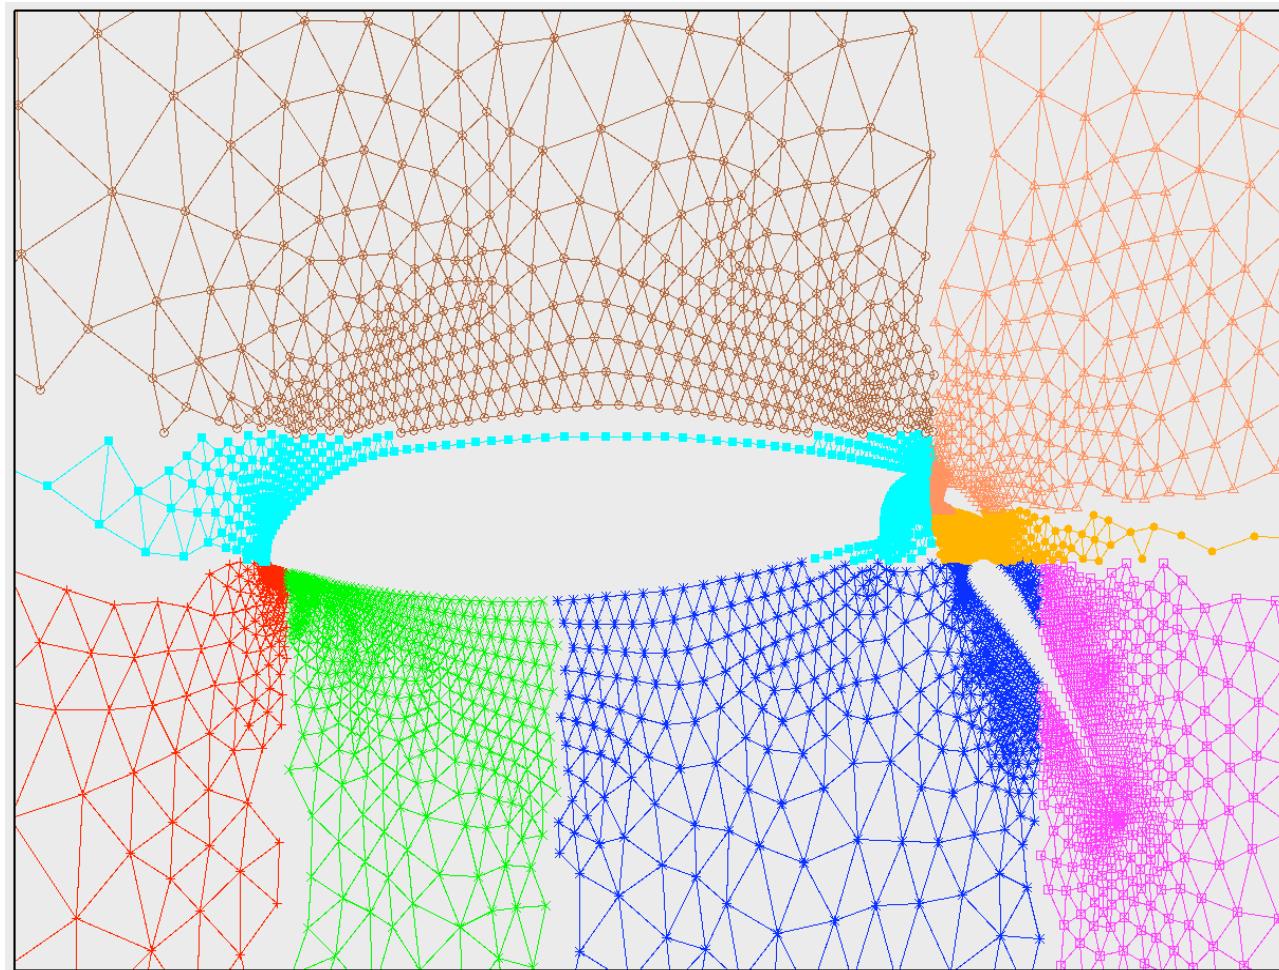
---

- Hammond airfoil mesh
- 2d mesh, triangular elements
  - 5K vertices
  - 13K edges
- Partition into 8 parts



# RCB

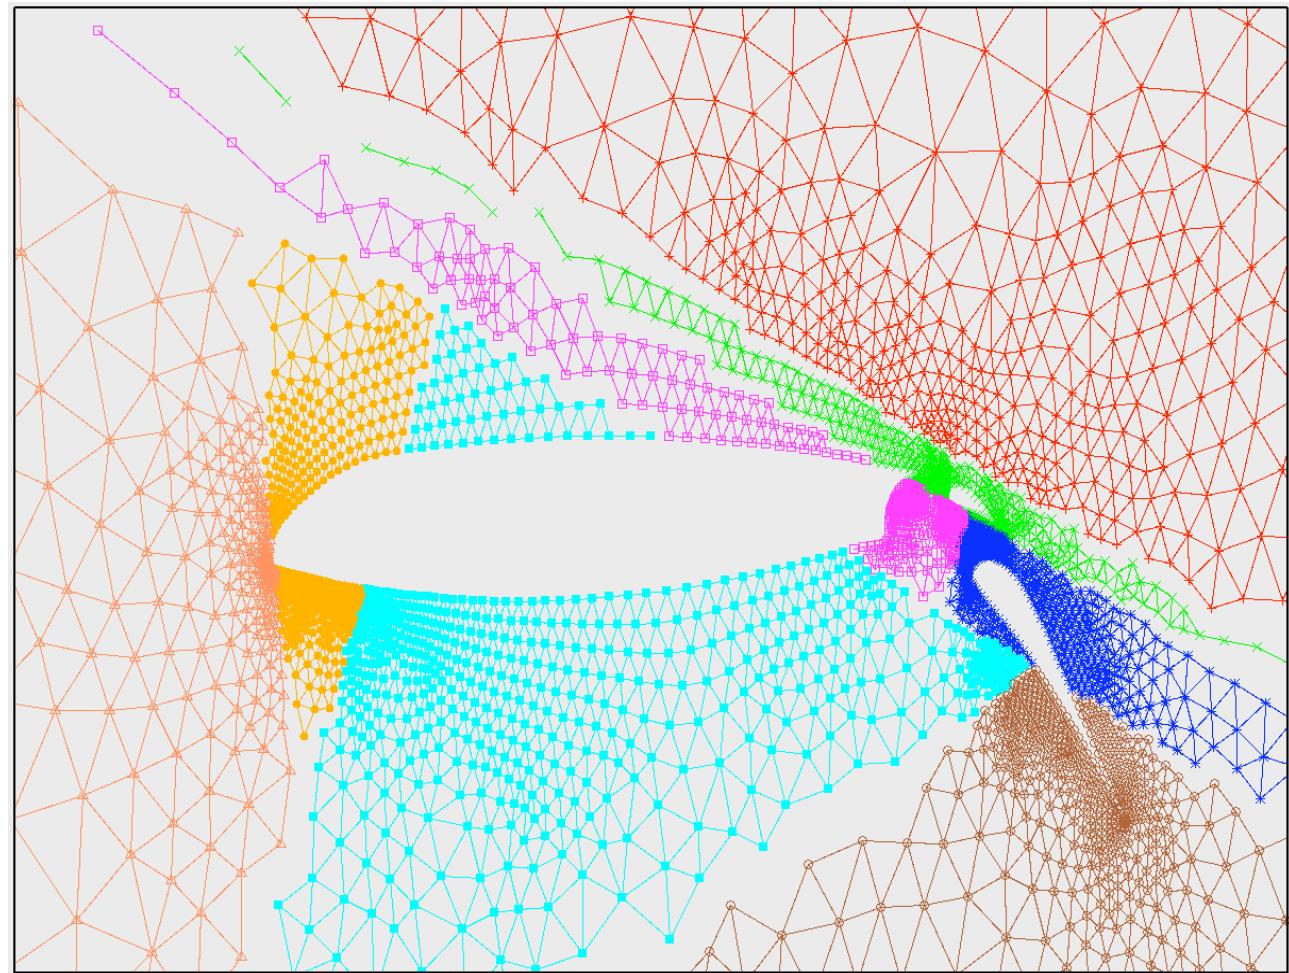
Total Volume:  
826  
Max #mesg:  
6





# RIB

Total volume:  
922  
Max #mesg: 5





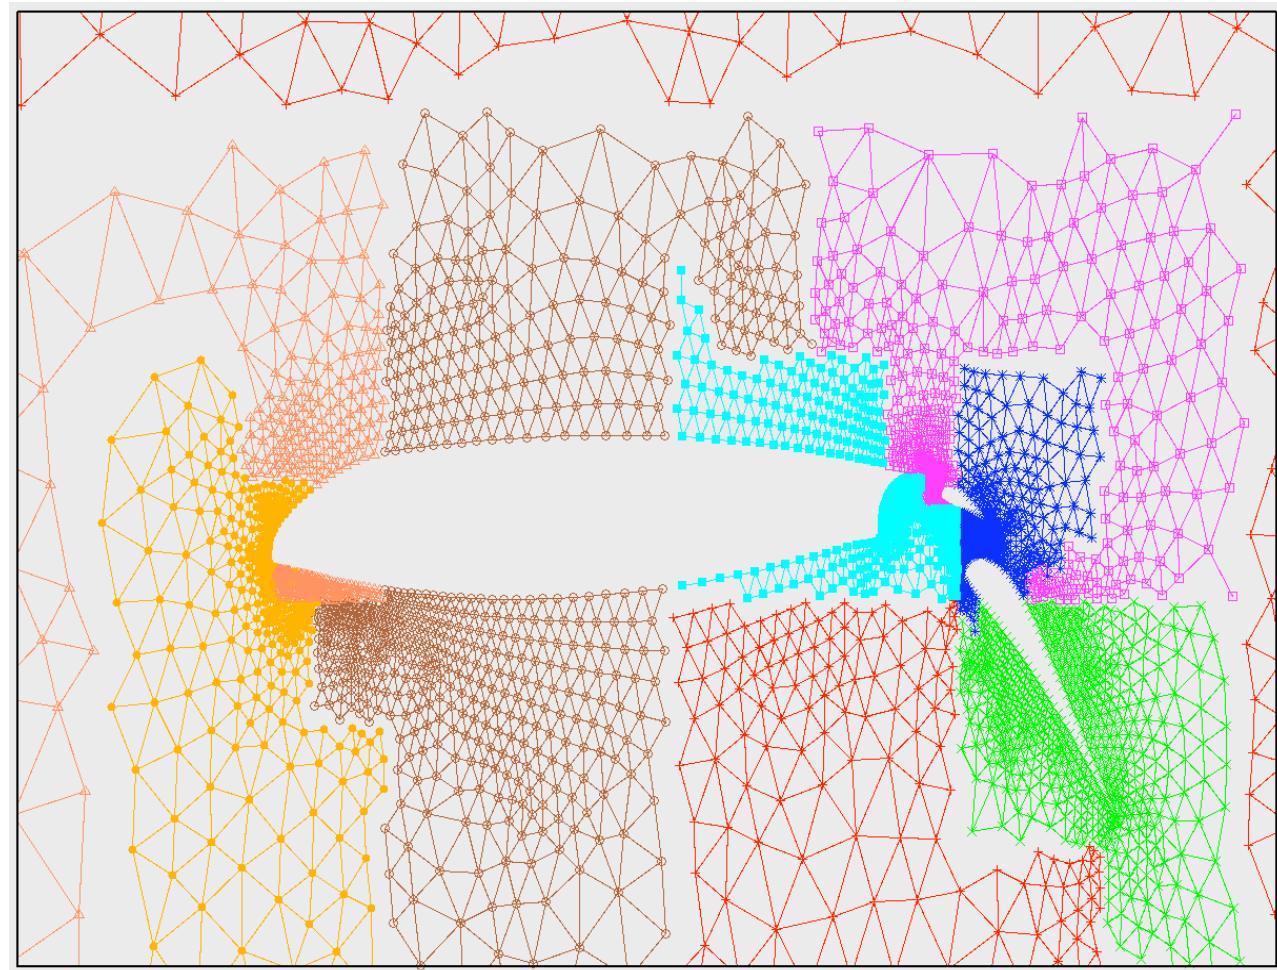
# HSFC

Total volume:

1000

Max #mesg:

6



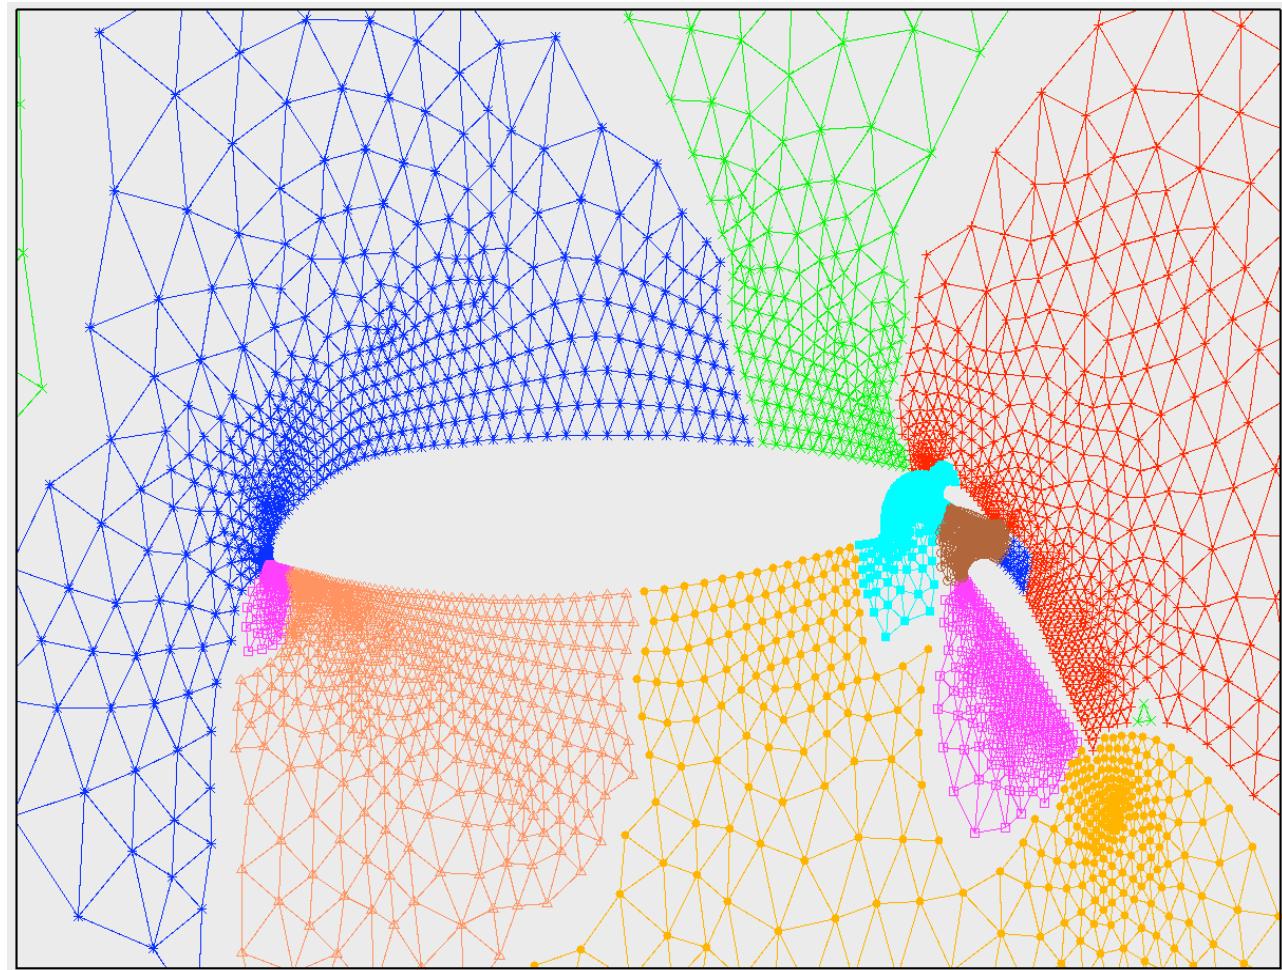


# Graph

Total volume:

472

Max #mesg: 5



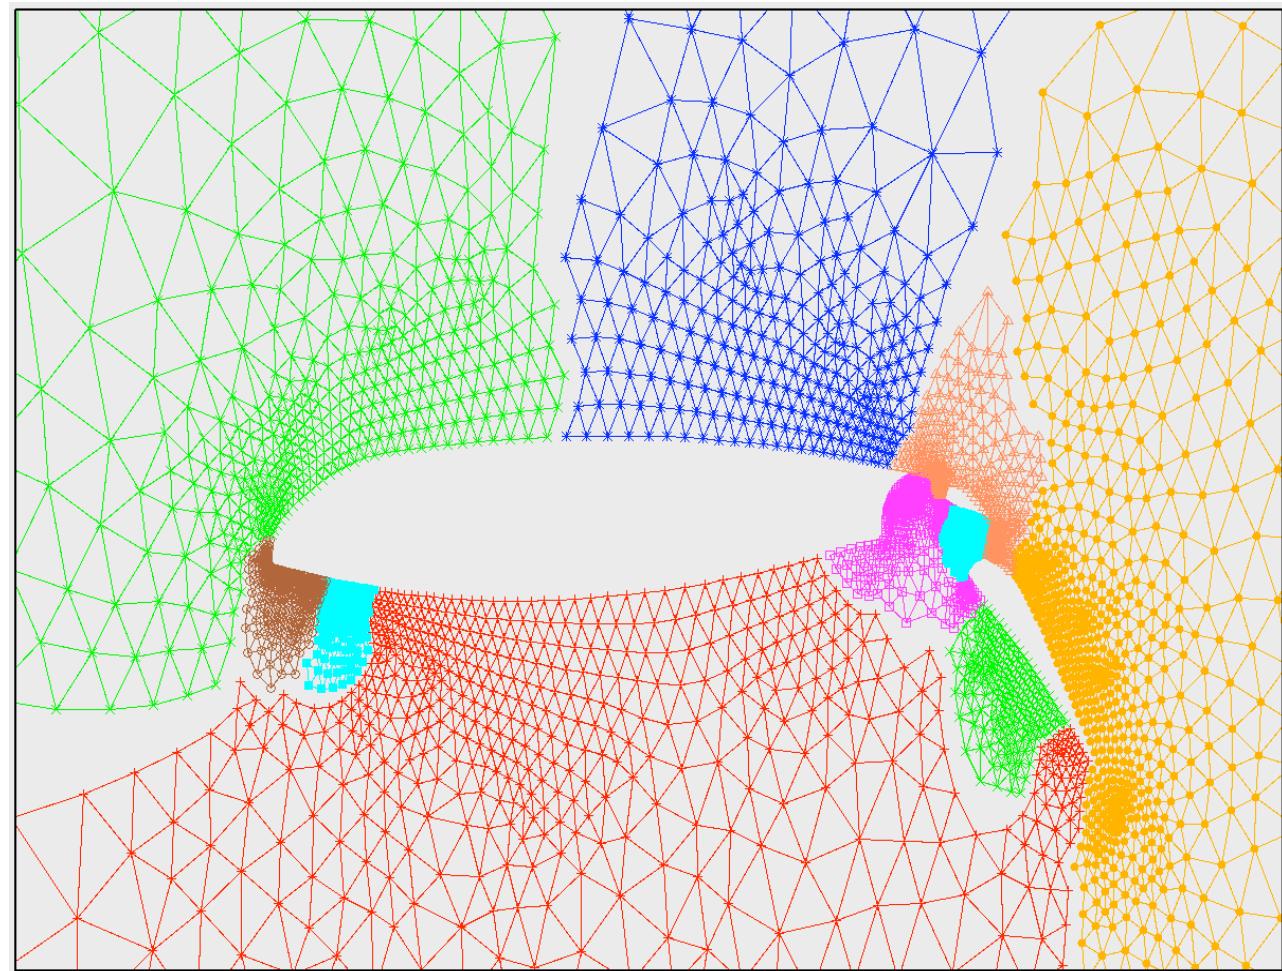


# Hypergraph

Total volume:

464

Max #mesg: 6





# Coffee Break!

---





# Software

---

- **Geometric partitioners**
  - Often embedded in application code;
    - Cannot easily be re-used.
- **Graph/hypergraph partitioners**
  - Multilevel partitioners are so complex they can take several man-years to implement.
  - Abstraction allows partitioners to be used across many applications.



# Software

---

- **1990s: Many graph partitioners**
  - Chaco (Sandia)
  - Metis/ParMetis (U. Minnesota)
  - Jostle/PJostle (U. Greenwich)
  - Scotch (U. Bordeaux)
  - Party (U. Paderborn)
- **Great advance at the time, but...**
  - Single algorithm is not best for all applications.
  - Interface requires application to build specific graph data structure.



# Our Approach: Zoltan Toolkit

---



- **Construct applications from smaller software “parts.”**
- “Tinker-toy parallel computing” -- B. Hendrickson
- **Toolkits include ...**
  - Services applications commonly need.
  - Support for wide range of applications.
  - Easy-to-use interfaces.
  - Data-structure neutral design.
- **Toolkits avoid ...**
  - Prescribed data structures
  - Heavy framework
  - Limited freedom for application developers.
- **Zoltan: Toolkit of Parallel Data Management Tools for Parallel, Unstructured Applications.**

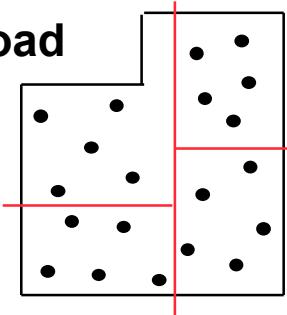


*Hasbro, Inc.*

# The Zoltan Toolkit

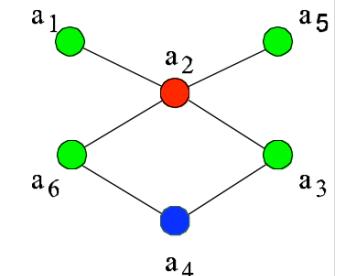
- Library of data management services for unstructured, dynamic and/or adaptive computations.

**Dynamic Load Balancing**

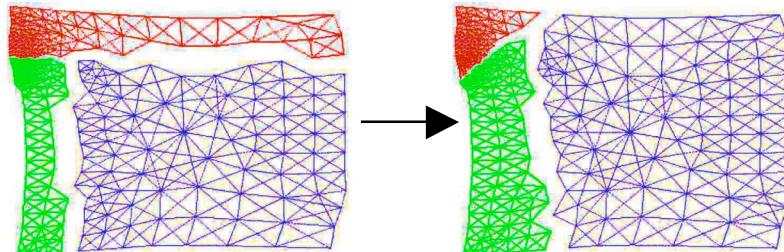


**Graph Coloring**

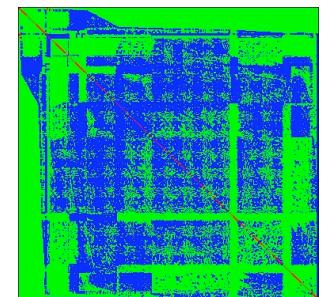
1	2	3	4	5	6
X	X				
X	X	X			
	X	X	X		
	X	X	X	X	
		X	X	X	X
			X	X	X
				X	X



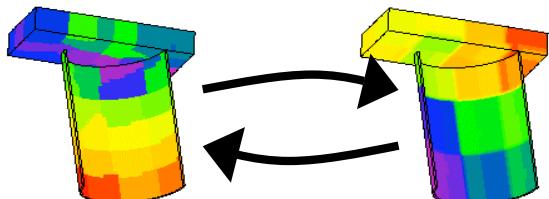
**Data Migration**



**Matrix Ordering**



**Unstructured Communication**



**Distributed Data Directories**

A	B	C	D	E	F	G	H	I
0	1	0	2	1	0	1	2	1

**Dynamic Memory Debugging**



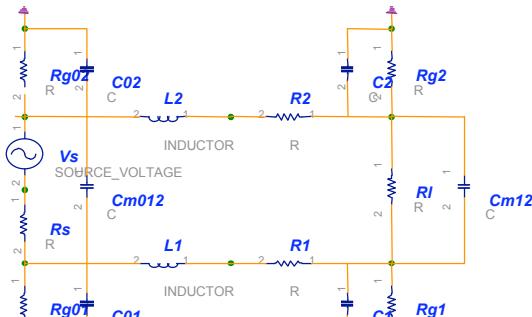


# Zoltan Supports Many Applications

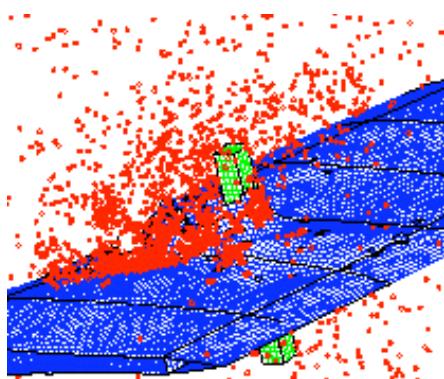
Slide 54



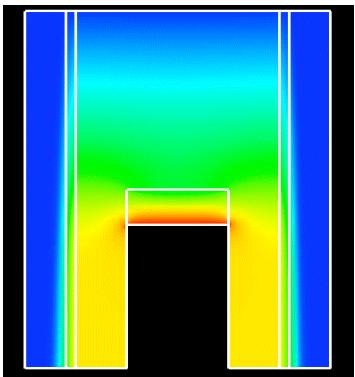
- Different applications, requirements, data structures.



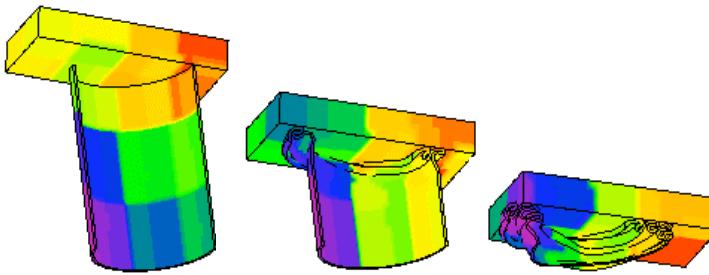
Parallel electronics networks



Particle methods



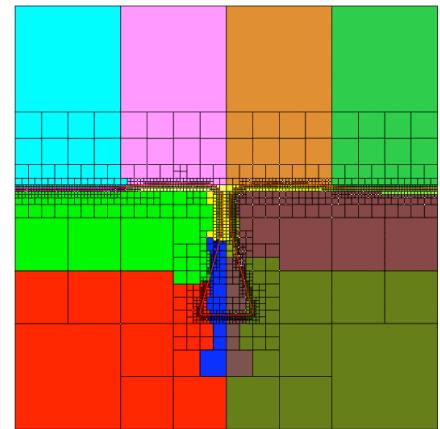
Multiphysics simulations



Crash simulations

$$\begin{matrix} & \text{A} & \\ \text{A} & \times & \text{b} \end{matrix} = \begin{matrix} & \text{b} & \end{matrix}$$

Linear solvers & preconditioners



Adaptive mesh refinement



# Zoltan Toolkit: Suite of Partitioners

---

Slide 55

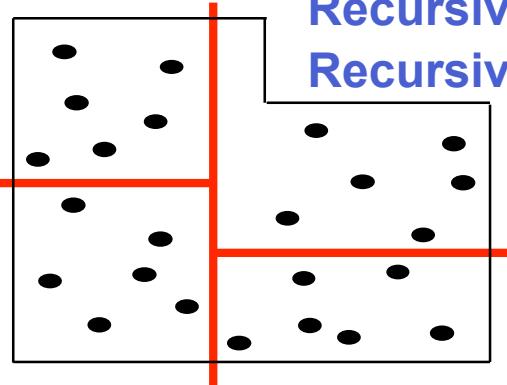


- No single partitioner works best for all applications.
  - Trade-offs:
    - Quality vs. speed.
    - Geometric locality vs. data dependencies.
    - High-data movement costs vs. tolerance for remapping.
- Application developers may not know which partitioner is best for application.
- Zoltan contains suite of partitioning methods.
  - Application changes only one parameter to switch methods.
    - `Zoltan_Set_Param(zz, "LB_METHOD", "new_method_name");`
  - Allows experimentation/comparisons to find most effective partitioner for application.



# Zoltan Toolkit: Suite of Partitioners

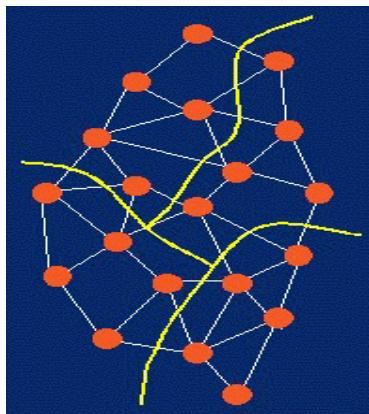
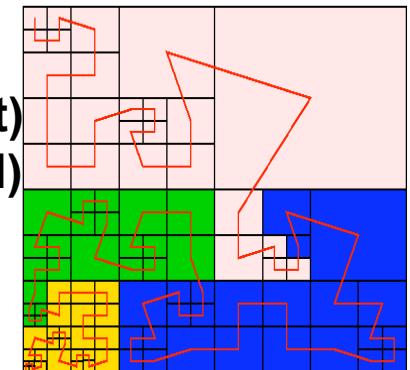
Slide 56



Recursive Coordinate Bisection (Berger, Bokhari)

Recursive Inertial Bisection (Taylor, Nour-Omid)

Space Filling Curves (Peano, Hilbert)  
Refinement-tree Partitioning (Mitchell)



Graph Partitioning

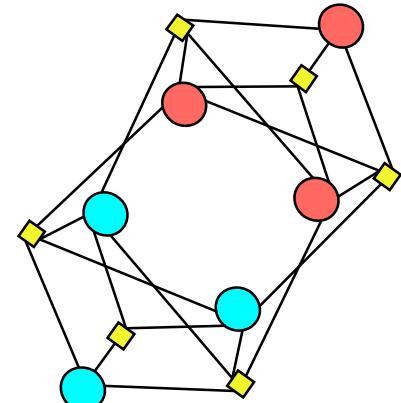
ParMETIS (Karypis, Schloegel, Kumar)

Jostle (Walshaw)

Hypergraph Partitioning & Repartitioning

(Catalyurek, Aykanat, Boman, Devine,  
Heaphy, Karypis, Bisseling)

PaToH (Catalyurek)



# Zoltan Interface Design

---

- Common interface to each class of partitioners.
- Partitioning method specified with user parameters.
  
- Data-structure neutral design.
  - Supports wide range of applications and data structures.
  - Imposes no restrictions on application's data structures.
  - Application does not have to build Zoltan's data structures.



# Zoltan Interface

---

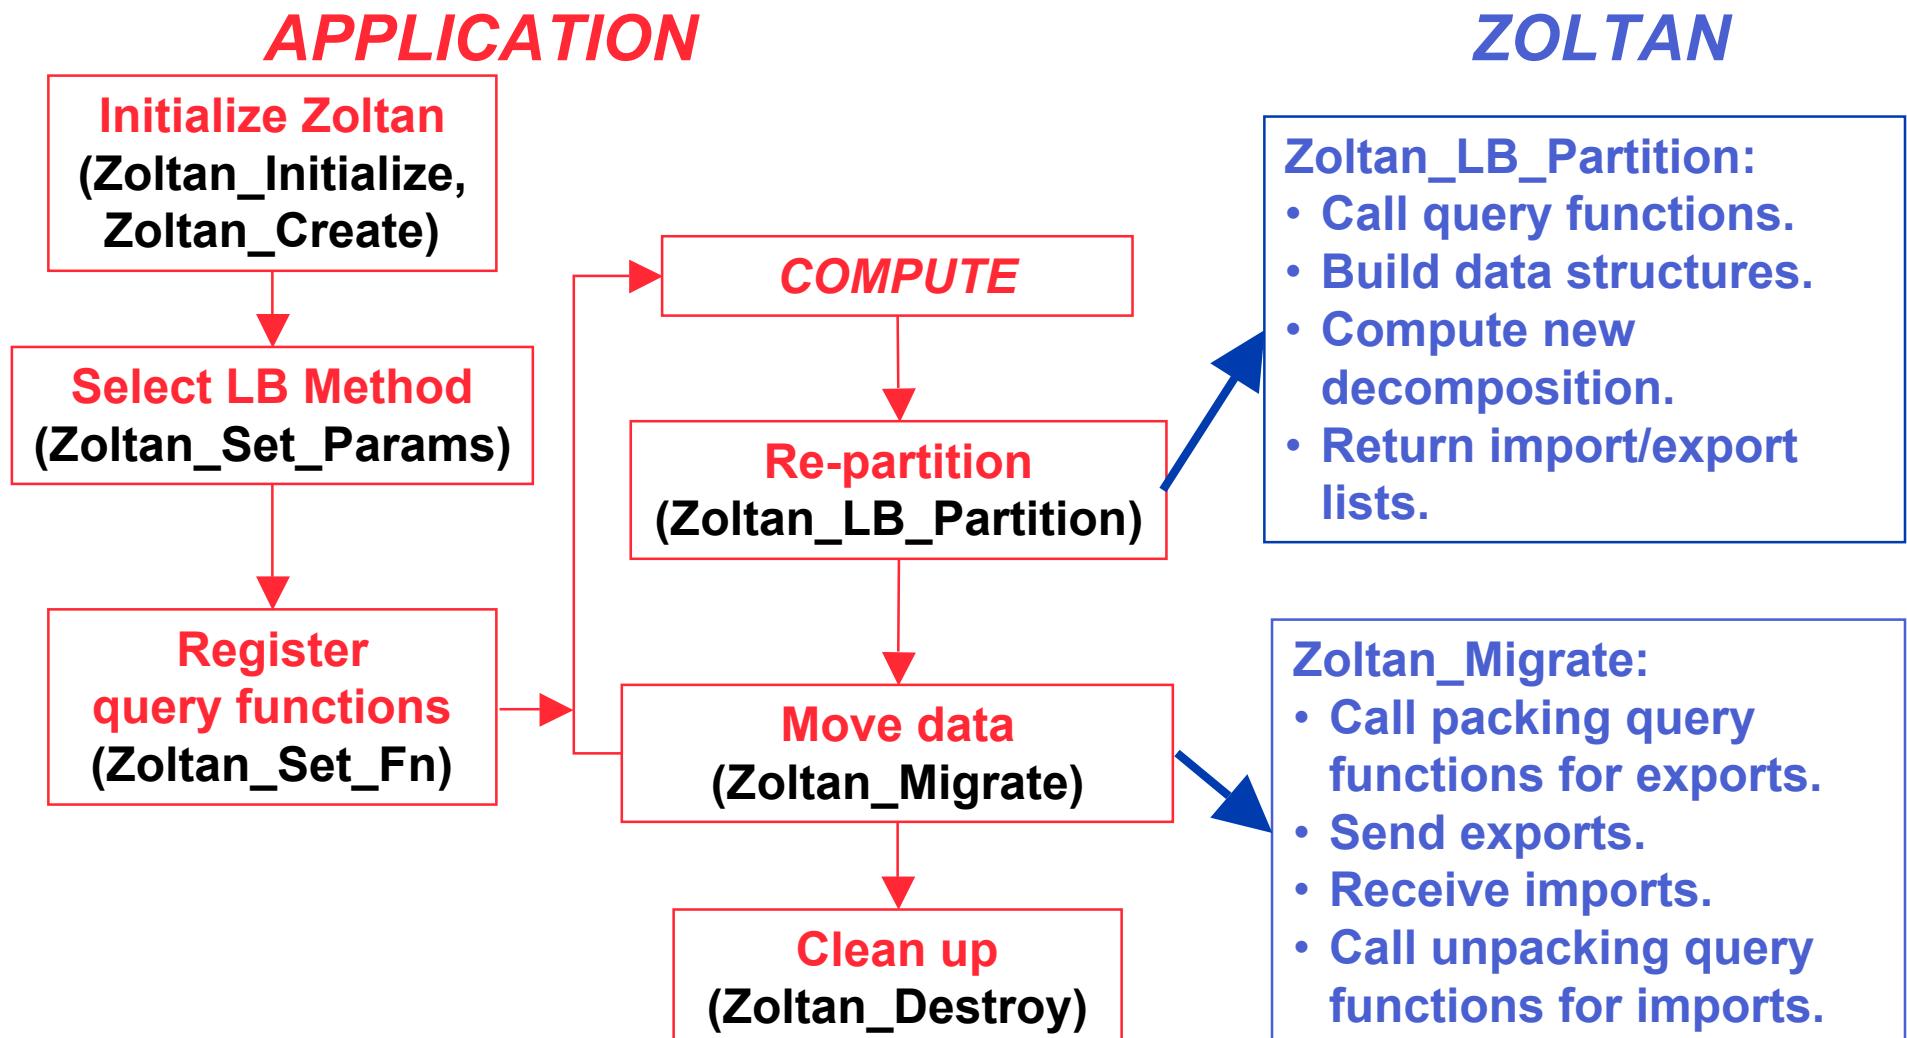
- Simple, easy-to-use interface.
  - Small number of callable Zoltan functions.
  - Callable from C, C++, Fortran.
- Requirement: Unique global IDs for objects to be partitioned. For example:
  - Global element number.
  - Global matrix row number.
  - (Processor number, local element number)
  - (Processor number, local particle number)

# Zoltan Application Interface

---

- Application interface:
  - Zoltan queries the application for needed info.
    - IDs of objects, coordinates, relationships to other objects.
  - Application provides simple functions to answer queries.
  - Small extra costs in memory and function-call overhead.
- Query mechanism supports...
  - Geometric algorithms
    - Queries for dimensions, coordinates, etc.
  - Hypergraph- and graph-based algorithms
    - Queries for edge lists, edge weights, etc.
  - Tree-based algorithms
    - Queries for parent/child relationships, etc.
- Once query functions are implemented, application can access all Zoltan functionality.
  - Can switch between algorithms by setting parameters.

# Zoltan Application Interface



# Zoltan Query Functions

## General Query Functions

<b>ZOLTAN_NUM_OBJ_FN</b>	Number of items on processor
<b>ZOLTAN_OBJ_LIST_FN</b>	List of item IDs and weights.

## Geometric Query Functions

<b>ZOLTAN_NUM_GEOM_FN</b>	Dimensionality of domain.
<b>ZOLTAN_GEOM_FN</b>	Coordinates of items.

## Hypergraph Query Functions

<b>ZOLTAN_HG_SIZE_CS_FN</b>	Number of hyperedge pins.
<b>ZOLTAN_HG_CS_FN</b>	List of hyperedge pins.
<b>ZOLTAN_HG_SIZE_EDGE_WTS_FN</b>	Number of hyperedge weights.
<b>ZOLTAN_HG_EDGE_WTS_FN</b>	List of hyperedge weights.

## Graph Query Functions

<b>ZOLTAN_NUM_EDGE_FN</b>	Number of graph edges.
<b>ZOLTAN_EDGE_LIST_FN</b>	List of graph edges.

# For geometric partitioning (RCB, RIB, HSFC), use ...

## General Query Functions

<b>ZOLTAN_NUM_OBJ_FN</b>	Number of items on processor
<b>ZOLTAN_OBJ_LIST_FN</b>	List of item IDs and weights.

## Geometric Query Functions

<b>ZOLTAN_NUM_GEOM_FN</b>	Dimensionality of domain.
<b>ZOLTAN_GEOM_FN</b>	Coordinates of items.

## Hypergraph Query Functions

<b>ZOLTAN_HG_SIZE_CS_FN</b>	Number of hyperedge pins.
<b>ZOLTAN_HG_CS_FN</b>	List of hyperedge pins.
<b>ZOLTAN_HG_SIZE_EDGE_WTS_FN</b>	Number of hyperedge weights.
<b>ZOLTAN_HG_EDGE_WTS_FN</b>	List of hyperedge weights.

## Graph Query Functions

<b>ZOLTAN_NUM_EDGE_FN</b>	Number of graph edges.
<b>ZOLTAN_EDGE_LIST_FN</b>	List of graph edges.



# For graph partitioning, coloring & ordering, use ...

Slide 63



## General Query Functions

<b>ZOLTAN_NUM_OBJ_FN</b>	Number of items on processor
<b>ZOLTAN_OBJ_LIST_FN</b>	List of item IDs and weights.

## Geometric Query Functions

<b>ZOLTAN_NUM_GEOM_FN</b>	Dimensionality of domain.
<b>ZOLTAN_GEOM_FN</b>	Coordinates of items.

## Hypergraph Query Functions

<b>ZOLTAN_HG_SIZE_CS_FN</b>	Number of hyperedge pins.
<b>ZOLTAN_HG_CS_FN</b>	List of hyperedge pins.
<b>ZOLTAN_HG_SIZE_EDGE_WTS_FN</b>	Number of hyperedge weights.
<b>ZOLTAN_HG_EDGE_WTS_FN</b>	List of hyperedge weights.

## Graph Query Functions

<b>ZOLTAN_NUM_EDGE_FN</b>	Number of graph edges.
<b>ZOLTAN_EDGE_LIST_FN</b>	List of graph edges.

# For hypergraph partitioning and repartitioning, use ...

<b>General Query Functions</b>	
<b>ZOLTAN_NUM_OBJ_FN</b>	Number of items on processor
<b>ZOLTAN_OBJ_LIST_FN</b>	List of item IDs and weights.
<b>Geometric Query Functions</b>	
<b>ZOLTAN_NUM_GEOM_FN</b>	Dimensionality of domain.
<b>ZOLTAN_GEOM_FN</b>	Coordinates of items.
<b>Hypergraph Query Functions</b>	
<b>ZOLTAN_HG_SIZE_CS_FN</b>	Number of hyperedge pins.
<b>ZOLTAN_HG_CS_FN</b>	List of hyperedge pins.
<b>ZOLTAN_HG_SIZE_EDGE_WTS_FN</b>	Number of hyperedge weights.
<b>ZOLTAN_HG_EDGE_WTS_FN</b>	List of hyperedge weights.
<b>Graph Query Functions</b>	
<b>ZOLTAN_NUM_EDGE_FN</b>	Number of graph edges.
<b>ZOLTAN_EDGE_LIST_FN</b>	List of graph edges.



# Or can use graph queries to build hypergraph.

Slide 65



General Query Functions	
ZOLTAN_NUM_OBJ_FN	Number of items on processor
ZOLTAN_OBJ_LIST_FN	List of item IDs and weights.
Geometric Query Functions	
ZOLTAN_NUM_GEOM_FN	Dimensionality of domain.
ZOLTAN_GEOM_FN	Coordinates of items.
Hypergraph Query Functions	
ZOLTAN_HG_SIZE_CS_FN	Number of hyperedge pins.
ZOLTAN_HG_CS_FN	List of hyperedge pins.
ZOLTAN_HG_SIZE_EDGE_WTS_FN	Number of hyperedge weights.
ZOLTAN_HG_EDGE_WTS_FN	List of hyperedge weights.
Graph Query Functions	
ZOLTAN_NUM_EDGE_FN	Number of graph edges.
ZOLTAN_EDGE_LIST_FN	List of graph edges.

# Using Zoltan in Your Application

---

1. Decide what your objects are.
  - Elements? Grid points? Matrix rows? Particles?
2. Decide which class of method to use  
(geometric/graph/hypergraph).
3. Download and build Zoltan.
4. Write required query functions for your application.
  - Required functions are listed with each method in Zoltan User's Guide.
5. Call Zoltan from your application.
6. #include “zoltan.h” in files calling Zoltan.
7. Compile; link application with libzoltan.a.
  - mpicc application.c -lzoltan



# Typical Applications

---

- **Unstructured meshes:**
  - Nodes, edges, and faces all need to be distributed.
  - Several choices:
    - Nodes are Zoltan objects (primal graph)
    - Faces are Zoltan objects (dual graph)
- **Sparse matrices:**
  - Partition rows or columns?
  - Balance rows or nonzeros?
- **Particle methods:**
  - Partition particles or cells weighted by particles?



# Zoltan: Getting Started

---

- Requirements:
  - C compiler
  - GNU Make (gmake)
  - MPI library (Message Passing Interface)
- Download Zoltan from Zoltan web site
  - <http://www.cs.sandia.gov/Zoltan>
  - Select “Download Zoltan” button.
  - Submit the registration form.
  - Choose the version you want;  
we suggest the latest version v3.0!
  - Downloaded file is `zoltan_distrib_v3.0.tar.gz`.

# Configuring and Building Zoltan

---

- Create and enter the Zoltan directory:
  - gunzip zoltan\_distrib\_v3.0.tar.gz
  - tar xf zoltan\_distrib\_v3.0.tar
  - cd Zoltan
- Configure and make Zoltan library
  - Not autotooled; uses manual configuration file.
  - “make zoltan” attempts a generic build;  
library libzoltan.a is in directory Obj\_generic.
  - To customize your build:
    - cd Utilities/Config; cp Config.linux Config.your\_system
    - Edit Config.your\_system
    - cd ../../
    - setenv ZOLTAN\_ARCH your\_system
    - make zoltan
    - Library libzoltan.a is in Obj\_your\_system



# Config file

---

```
DEFS          =
RANLIB        = ranlib
AR            = ar r

CC            = mpicc -Wall
CPPC          = mpic++
INCLUDE_PATH  =
DBG_FLAGS     = -g
OPT_FLAGS     = -O
CFLAGS         = $(DBG_FLAGS)

F90           = mpif90
LOCAL_F90     = f90
F90CFLAGS    = -DFMANGLE=UNDERSCORE -DNO_MPI2
FFLAGS        =
SPPR_HEAD     = spprinc.most
F90_MODULE_PREFIX = -I
FARG          = farg_typical

MPI_LIB        =
MPI_LIBPATH   =

PARMETIS_LIBPATH = -L/Users/kddevin/code/ParMETIS3_1
PARMETIS_INCPATH = -I/Users/kddevin/code/ParMETIS3_1
#PATOH_LIBPATH  = -L/Users/kddevin/code/PaToH
#PATOH_INCPATH = -I/Users/kddevin/code/PaToH
```



# Simple Example

---

- **Zoltan/examples/C/zoltanSimple.c**
- **Application data structure:**
  - **int MyNumPts;**
    - Number of points on processor.
  - **int \*Gids;**
    - array of Global ID numbers of points on processor.
  - **float \*Pts;**
    - Array of 3D coordinates of points on processor (in same order as Gids array).



# Example zoltanSimple.c: Initialization

---

Slide 72



```
/* Initialize MPI */
MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &me);
MPI_Comm_size(MPI_COMM_WORLD, &nprocs);

/*
** Initialize application data. In this example,
** create a small test mesh and divide it across processors
*/

exSetDivisions(32);      /* rectilinear mesh is div X div X div */

MyNumPts = exInitializePoints(&Pts, &Gids, me, nprocs);

/* Initialize Zoltan */
rc = Zoltan_Initialize(argc, argv, &ver);

if (rc != ZOLTAN_OK){
    printf("sorry...\n");
    free(Pts); free(Gids);
    exit(0);
}
```



# Example zoltanSimple.c: Prepare for Partitioning

---

Slide 73



```
/* Allocate and initialize memory for Zoltan structure */
zz = Zoltan_Create(MPI_COMM_WORLD);

/* Set general parameters */
Zoltan_Set_Param(zz, "DEBUG_LEVEL", "0");
Zoltan_Set_Param(zz, "LB_METHOD", "RCB");
Zoltan_Set_Param(zz, "NUM_GID_ENTRIES", "1");
Zoltan_Set_Param(zz, "NUM_LID_ENTRIES", "1");
Zoltan_Set_Param(zz, "RETURN_LISTS", "ALL");

/* Set RCB parameters */
Zoltan_Set_Param(zz, "KEEP_CUTS", "1");
Zoltan_Set_Param(zz, "RCB_OUTPUT_LEVEL", "0");
Zoltan_Set_Param(zz, "RCB_RECTILINEAR_BLOCKS", "1");

/* Register call-back query functions. */
Zoltan_Set_Num_Obj_Fn(zz, exGetNumberOfAssignedObjects, NULL);
Zoltan_Set_Obj_List_Fn(zz, exGetObjectList, NULL);
Zoltan_Set_Num_Geom_Fn(zz, exGetObjectSize, NULL);
Zoltan_Set_Geom_Multi_Fn(zz, exGetObject, NULL);
```



# Example zoltanSimple.c: Partitioning

Slide 74



**Sandia  
National  
Laboratories**

Zoltan computes the **difference** ( $\Delta$ ) from current distribution  
Choose between:

- a) Import lists (data to import **from** other procs)
  - b) Export lists (data to export **to** other procs)
  - c) Both (the default)



# Example zoltanSimple.c: Use the Partition

---

Slide 75



```
/* Process partitioning results;
** in this case, print information;
** in a "real" application, migrate data here.
*/
if (!rc){
    exPrintGlobalResult("Recursive Coordinate Bisection",
                        nprocs, me,
                        MyNumPts, numImport, numExport, changes);
}
else{
    free(Pts);
    free(Gids);
    Zoltan_Destroy(&zz);
    MPI_Finalize();
    exit(0);
}
```



# Example zoltanSimple.c: Cleanup

---

Slide 76



```
/* Free Zoltan memory allocated by Zoltan_LB_Partition. */
Zoltan_LB_Free_Part(&importGlobalGids, &importLocalGids,
                     &importProcs, &importToPart);
Zoltan_LB_Free_Part(&exportGlobalGids, &exportLocalGids,
                     &exportProcs, &exportToPart);

/* Free Zoltan memory allocated by Zoltan_Create. */
Zoltan_Destroy(&zz);

/* Free Application memory */
free(Pts); free(Gids);

/******************
** all done *****
****************/

MPI_Finalize();
```



# Example zoltanSimple.c: ZOLTAN\_OBJ\_LIST\_FN

Slide 77



```
void exGetObjectList(void *userDefinedData,
                     int numGlobalIds, int numLocalIds,
                     ZOLTAN_ID_PTR gids, ZOLTAN_ID_PTR lids,
                     int wgt_dim, float *obj_wgts,
                     int *err)

{
/* ZOLTAN_OBJ_LIST_FN callback function.
** Returns list of objects owned by this processor.
** lids[i] = local index of object in array.
*/
    int i;

    for (i=0; i<NumPoints; i++)
    {
        gids[i] = GlobalIds[i];
        lids[i] = i;
    }

    *err = 0;

    return;
}
```



# Example zoltanSimple.c:

## ZOLTAN\_GEOOM\_MULTI\_FN

Slide 78



```
void exGetObjectCoords(void *userDefinedData,
                      int numGlobalIds, int numLocalIds, int numObjs,
                      ZOLTAN_ID_PTR gids, ZOLTAN_ID_PTR lids,
                      int numDim, double *pts, int *err)
{
    /* ZOLTAN_GEOOM_MULTI_FN callback.
     ** Returns coordinates of objects listed in gids and lids.
    */
    int i, id, id3, next = 0;
    if (numDim != 3) {
        *err = 1; return;
    }
    for (i=0; i<numObjs; i++){
        id = lids[i];
        if ((id < 0) || (id >= NumPoints)) {
            *err = 1; return;
        }
        id3 = lids[i] * 3;
        pts[next++] = (double)(Points[id3]);
        pts[next++] = (double)(Points[id3 + 1]);
        pts[next++] = (double)(Points[id3 + 2]);
    }
}
```

# Example Graph Callbacks

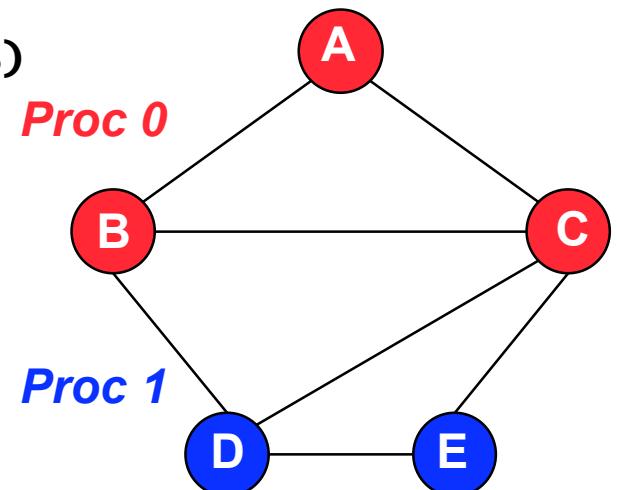
```
void ZOLTAN_NUM_EDGES_MULTI_FN(void *data,
    int num_gid_entries, int num_lid_entries,
    int num_obj, ZOLTAN_ID_PTR global_id, ZOLTAN_ID_PTR local_id,
    int *num_edges, int *ierr);
```

Proc 0 Input from Zoltan:

```
num_obj = 3
global_id = {A,C,B}
local_id = {0,1,2}
```

Output from Application on Proc 0:

```
num_edges = {2,4,3}
            (i.e., degrees of vertices A, C, B)
ierr = ZOLTAN_OK
```



# Example Graph Callbacks

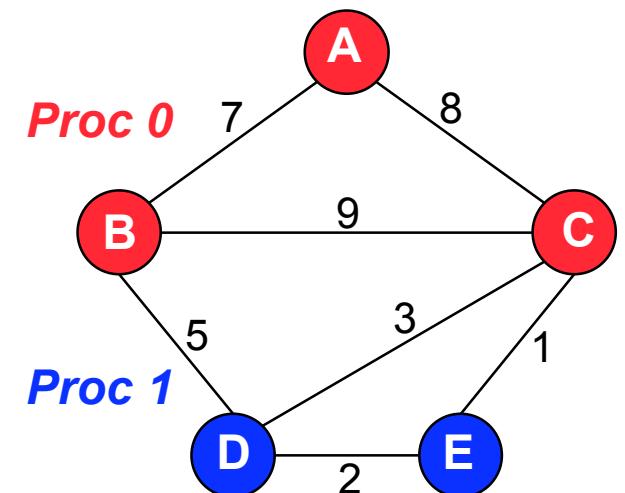
```
void ZOLTAN_EDGE_LIST_MULTI_FN(void *data,
    int num_gid_entries, int num_lid_entries,
    int num_obj, ZOLTAN_ID_PTR global_ids, ZOLTAN_ID_PTR local_ids,
    int *num_edges,
    ZOLTAN_ID_PTR nbor_global_id, int *nbor_procs,
    int wdim, float *nbor_ewgts,
    int *ierr);
```

**Proc 0 Input from Zoltan:**

**num\_obj** = 3  
**global\_ids** = {A, C, B}  
**local\_ids** = {0, 1, 2}  
**num\_edges** = {2, 4, 3}  
**wdim** = 0 or EDGE\_WEIGHT\_DIM parameter value

**Output from Application on Proc 0:**

**nbor\_global\_id** = {B, C, A, B, E, D, A, C, D}  
**nbor\_procs** = {0, 0, 0, 0, 1, 1, 0, 0, 1}  
**nbor\_ewgts** = if **wdim** then  
{7, 8, 8, 9, 1, 3, 7, 9, 5}  
**ierr** = ZOLTAN\_OK



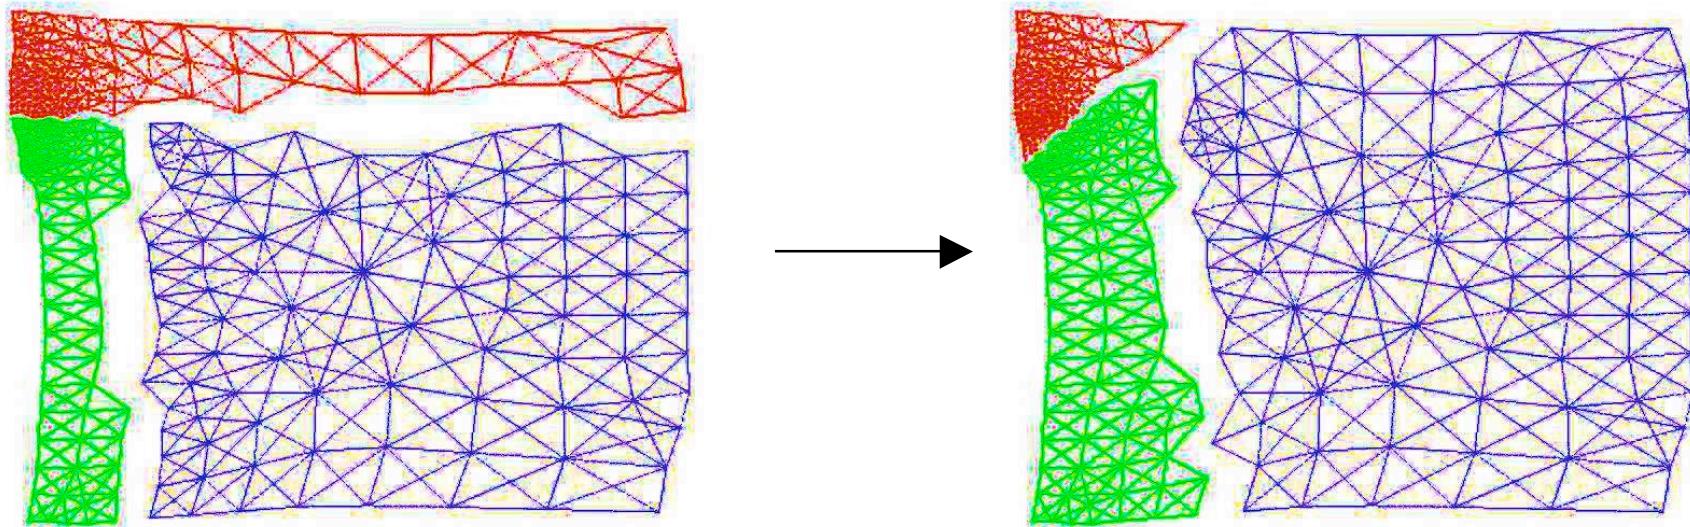


# More Details on Query Functions

- **void\* data pointer** allows user data structures to be used in all query functions.
    - To use, cast the pointer to the application data type.
  - Local IDs provided by application are returned by Zoltan to simplify access of application data.
    - E.g. Indices into local arrays of coordinates.
  - **ZOLTAN\_ID\_PTR** is pointer to array of unsigned integers, allowing IDs to be more than one integer long.
    - E.g., (processor number, local element number) pair.
    - **numGlobalIds** and **numLocalIds** are lengths of each ID.
  - All memory for query-function arguments is allocated in Zoltan.

# Zoltan Data Migration Tools

- After partition is computed, data must be moved to new decomposition.
  - Depends strongly on application data structures.
  - Complicated communication patterns.
- Zoltan can help!
  - Application supplies query functions to pack/unpack data.
  - Zoltan does all communication to new processors.





## Data Migration Tools

---

- Required migration query functions:
  - **ZOLTAN\_OBJ\_SIZE\_MULTI\_FN:**
    - Returns size of data (in bytes) for each object to be exported to a new processor.
  - **ZOLTAN\_PACK\_MULTI\_FN:**
    - Remove data from application data structure on old processor;
    - Copy data to Zoltan communication buffer.
  - **ZOLTAN\_UNPACK\_MULTI\_FN:**
    - Copy data from Zoltan communication buffer into data structure on new processor.
- ```
int Zoltan_Migrate(struct Zoltan_Struct *zz,
                    int num_import, ZOLTAN_ID_PTR import_global_ids,
                    ZOLTAN_ID_PTR import_local_ids, int *import_procs,
                    int *import_to_part,
                    int num_export, ZOLTAN_ID_PTR export_global_ids,
                    ZOLTAN_ID_PTR export_local_ids, int *export_procs,
                    int *export_to_part);
```

# Other Zoltan Functionality

---

- Tools needed when doing dynamic load balancing:
  - Unstructured Communication Primitives
  - Distributed Data Directories
- Tools closely related to graph partitioning:
  - Graph coloring
  - Matrix ordering
  - These tools use the same query functions as graph partitioners.
- All functionality described in Zoltan User's Guide.
  - [http://www.cs.sandia.gov/Zoltan/ug\\_html/ug.html](http://www.cs.sandia.gov/Zoltan/ug_html/ug.html)

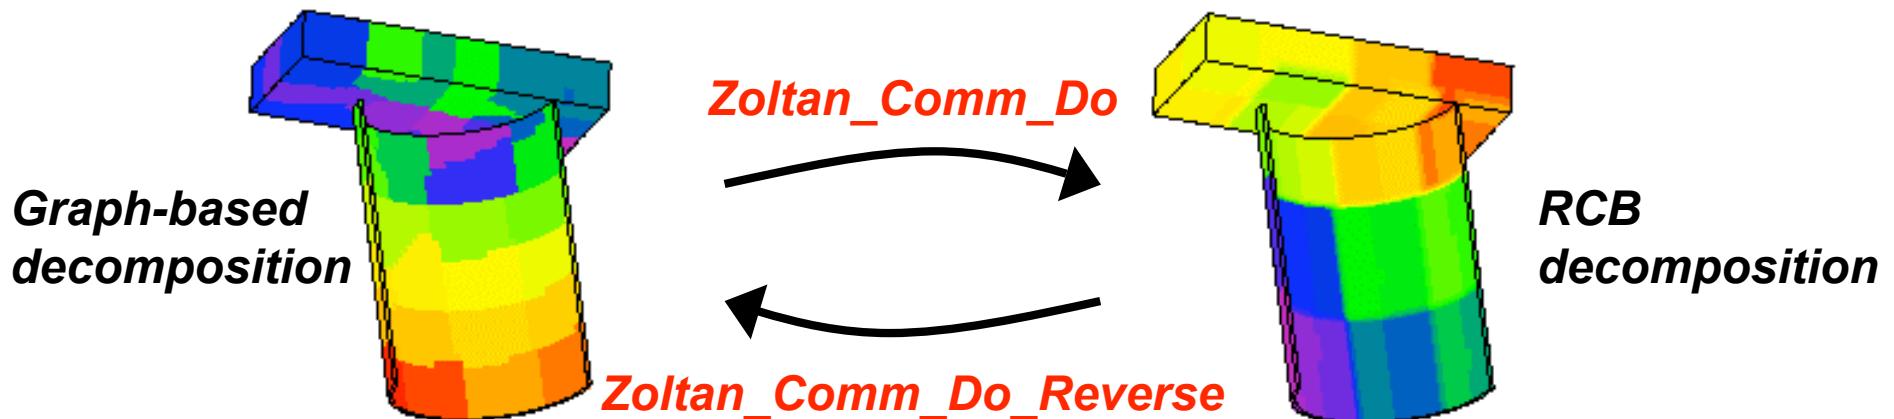


# Zoltan Unstructured Communication Package

Slide 85



- Simple primitives for efficient irregular communication.
  - Zoltan\_Comm\_Create: Generates communication plan.
    - Processors and amount of data to send and receive.
  - Zoltan\_Comm\_Do: Send data using plan.
    - Can reuse plan. (Same plan, different data.)
  - Zoltan\_Comm\_Do\_Reverse: Inverse communication.
- Used for most communication in Zoltan.





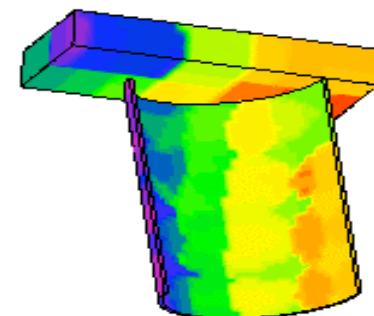
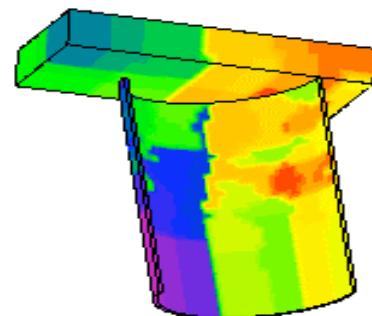
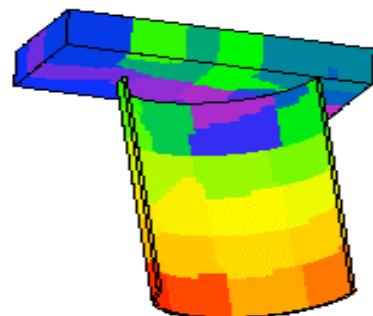
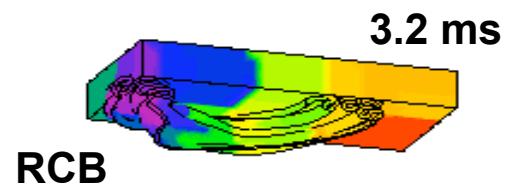
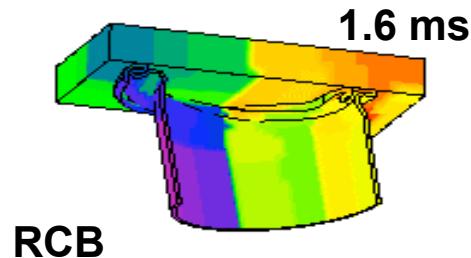
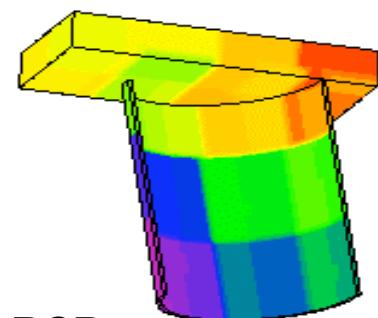
# Example Application: Crash Simulations

Slide 86



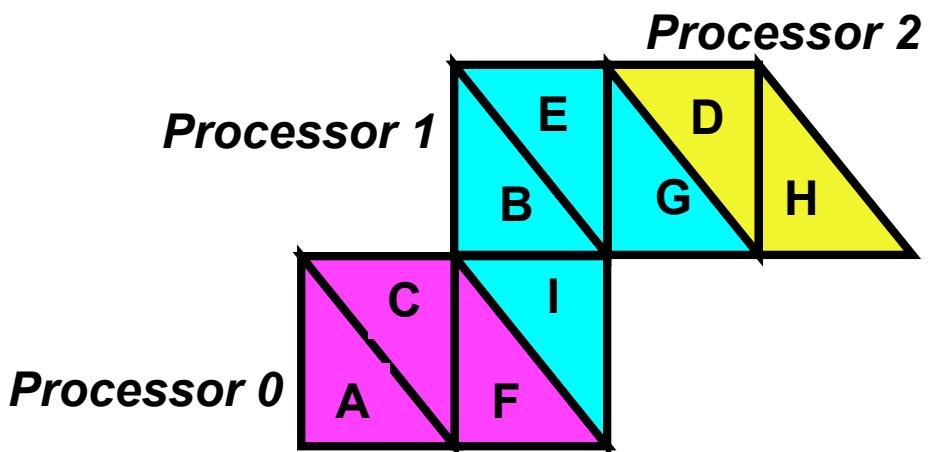
- **Multiphase simulation:**

- Graph-based decomposition of elements for finite element calculation.
- Dynamic geometric decomposition of surfaces for contact detection.
- Migration tools and Unstructured Communication package map between decompositions.



# Zoltan Distributed Data Directory

- Helps applications locate off-processor data.
- Rendezvous algorithm (Pinar, 2001).
  - Directory distributed in known way (hashing) across processors.
  - Requests for object location sent to processor storing the object's directory entry.



|                          |   |   |   |
|--------------------------|---|---|---|
| <i>Directory Index →</i> | A | B | C |
| <i>Location →</i>        | 0 | 1 | 0 |

Processor 0

|   |   |   |
|---|---|---|
| D | E | F |
| 2 | 1 | 0 |

Processor 1

|   |   |   |
|---|---|---|
| G | H | I |
| 1 | 2 | 1 |

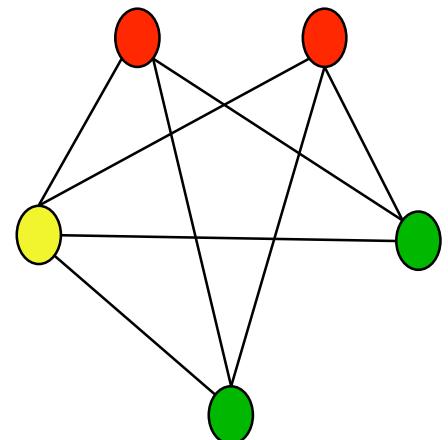
Processor 2



# Zoltan Graph Coloring

---

- Parallel distance-1 and distance-2 graph coloring.
- Graph built using same application interface and code as graph partitioners.
- Generic coloring interface; easy to add new coloring algorithms.
- Implemented algorithms due to Bozdag, Catalyurek, Gebremedhin, Manne, Boman, 2005.

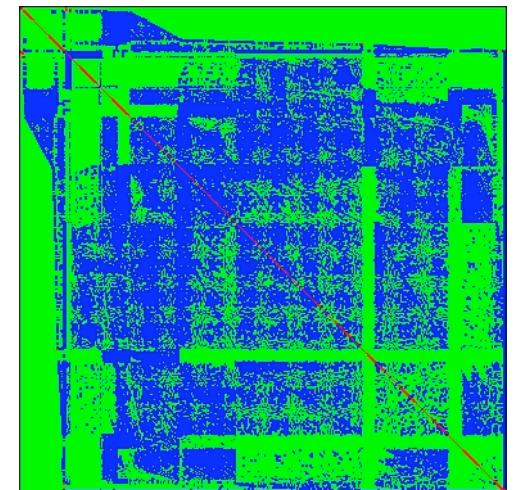




# Zoltan Matrix Ordering Interface

---

- Produce **fill-reducing ordering for sparse matrix factorization.**
- Graph built using same application interface and code as graph partitioners.
- Generic ordering interface; easy to add new ordering algorithms.
- Specific interface to ordering methods in ParMETIS (Karypis, et al., U. Minnesota).



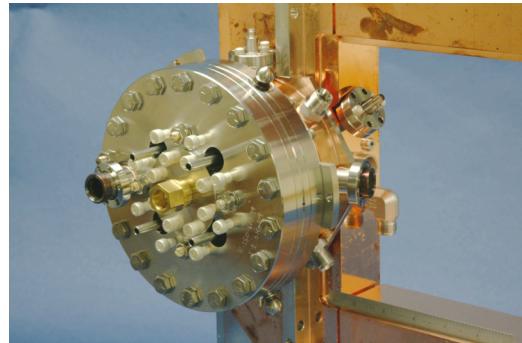


# Performance Results

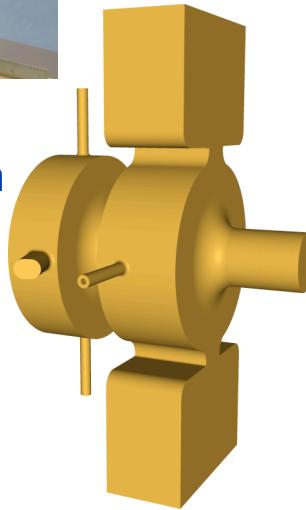
---

- Experiments on Sandia's Thunderbird cluster.
  - Dual 3.6 GHz Intel EM64T processors with 6 GB RAM.
  - Infiniband network.
- Compare RCB, graph and hypergraph methods.
- Measure ...
  - Amount of communication induced by the partition.
  - Partitioning time.

# Test Data



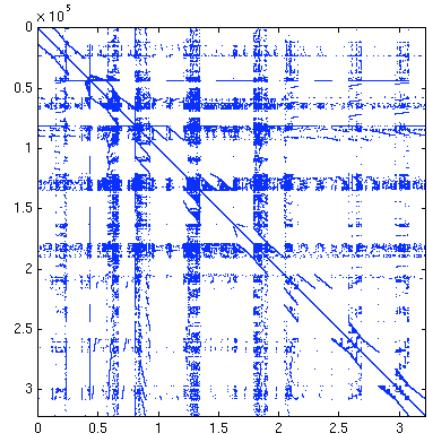
**SLAC \*LCLS**  
**Radio Frequency Gun**  
**6.0M x 6.0M**  
**23.4M nonzeros**



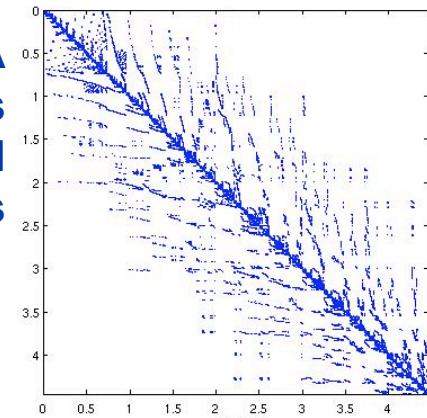
**SLAC Linear Accelerator**  
**2.9M x 2.9M**  
**11.4M nonzeros**



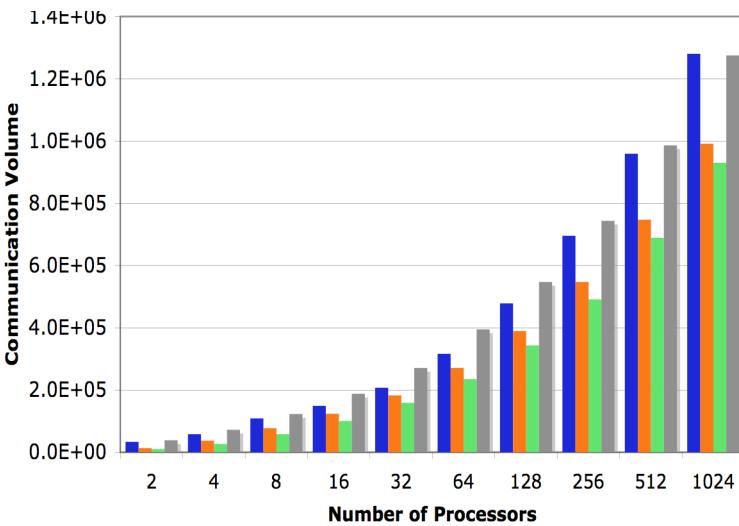
**Xyce 680K ASIC Stripped  
Circuit Simulation**  
**680K x 680K**  
**2.3M nonzeros**



**Cage15 DNA  
Electrophoresis**  
**5.1M x 5.1M**  
**99M nonzeros**

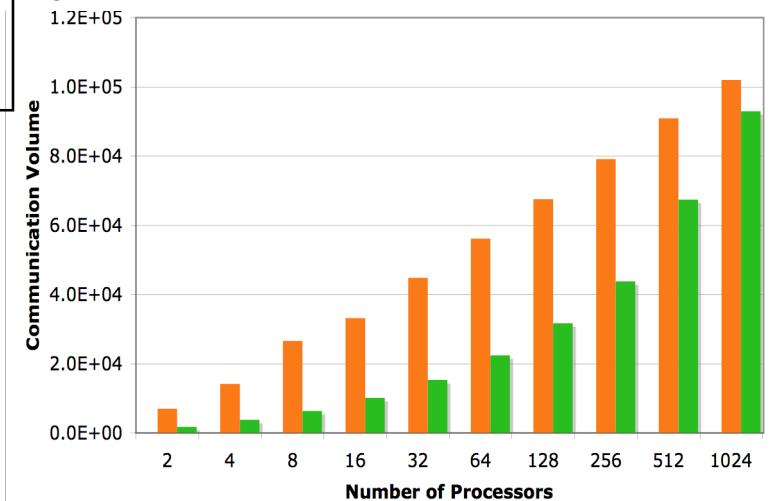
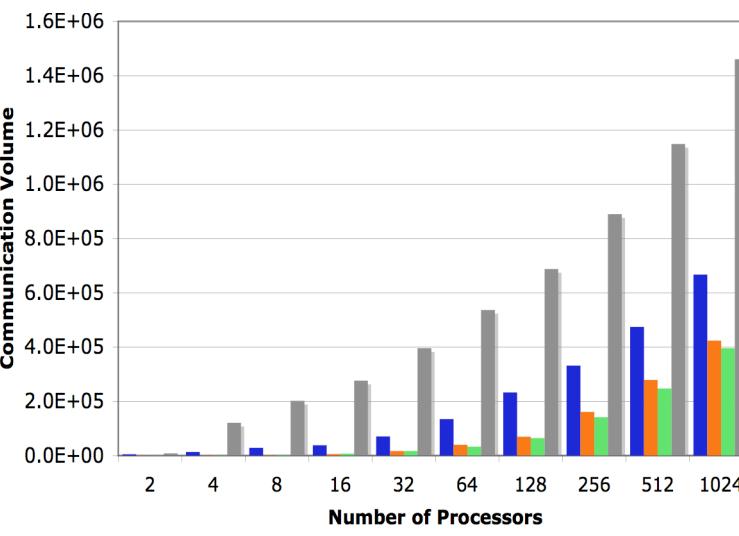
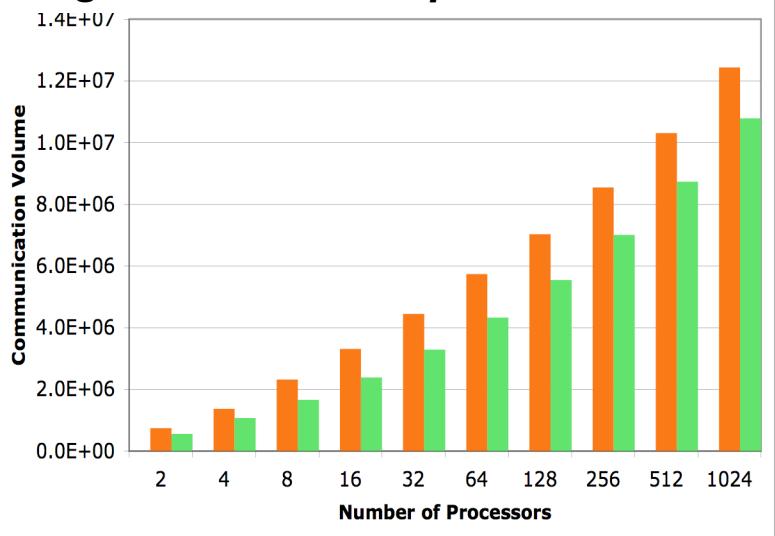


# Communication Volume: Lower is Better

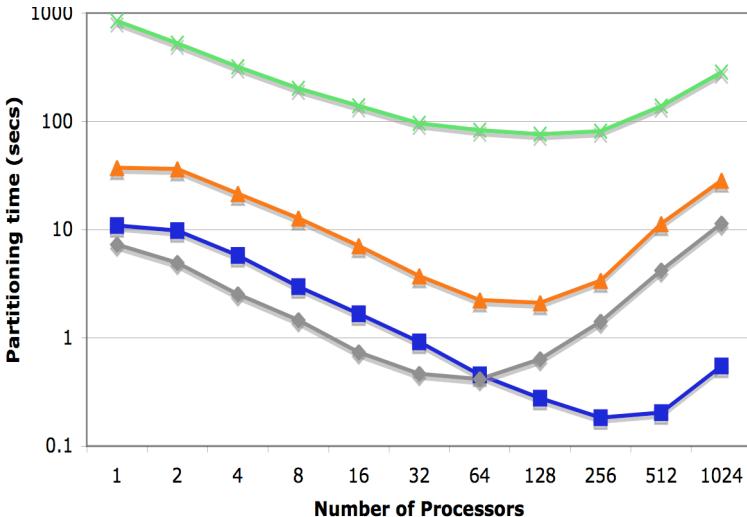
**SLAC 6.0M LCLS**

*Number of parts  
= number of  
processors.*

**RCB**  
**Graph**  
**Hypergraph**  
**HSFC**

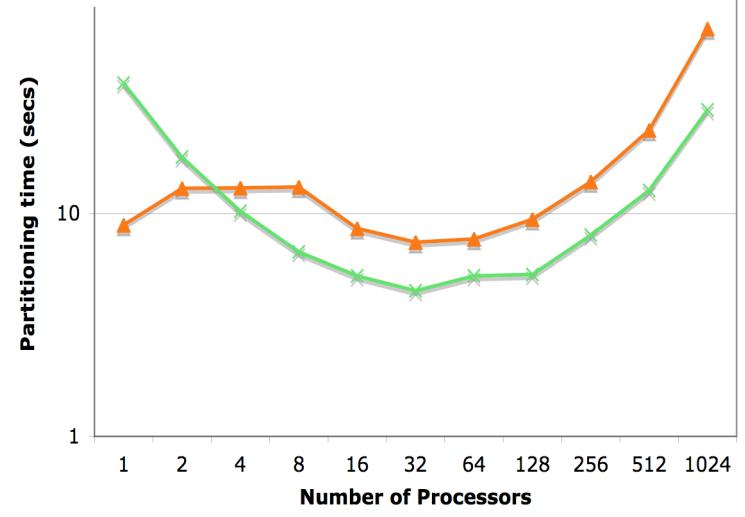
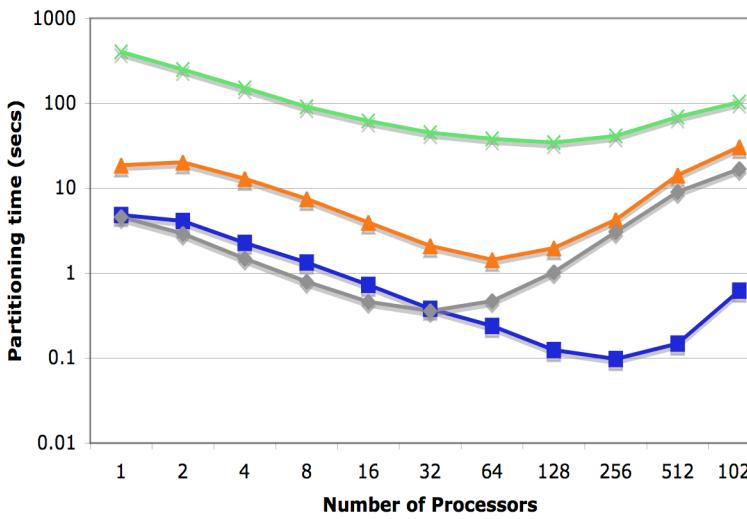
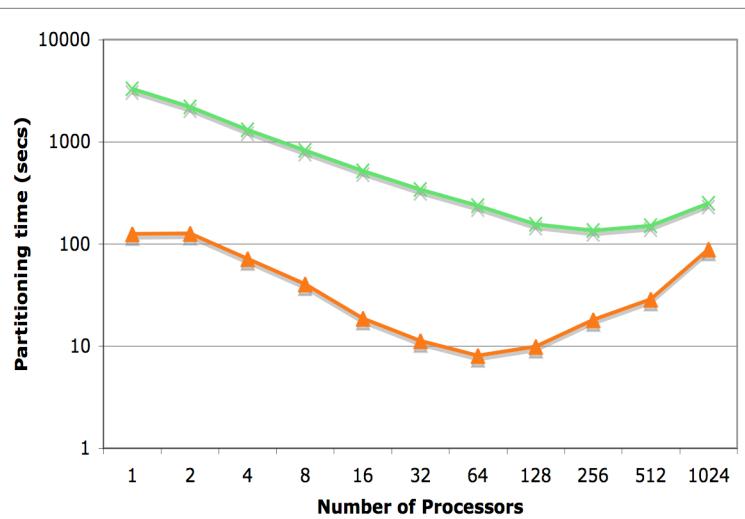
**Xyce 680K circuit****SLAC 2.9M Linear Accelerator****Cage15 5.1M electrophoresis**

# Partitioning Time: Lower is better

**SLAC 6.0M LCLS**

1024 parts.  
Varying number  
of processors.

- █ RCB
- █ Graph
- █ Hypergraph
- █ HSFC

**Xyce 680K circuit****SLAC 2.9M Linear Accelerator****Cage15 5.1M electrophoresis**

# Repartitioning Experiments

---

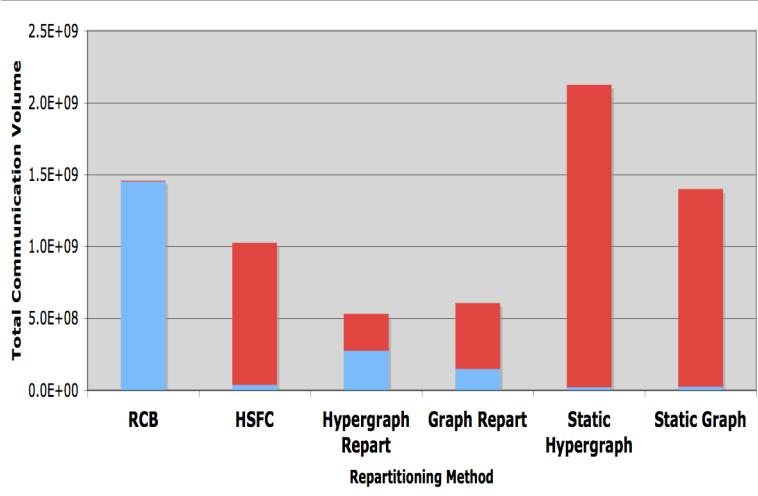
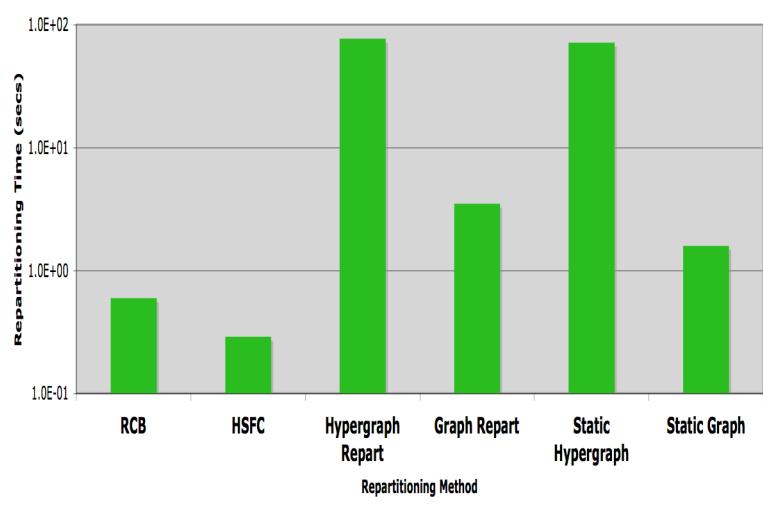
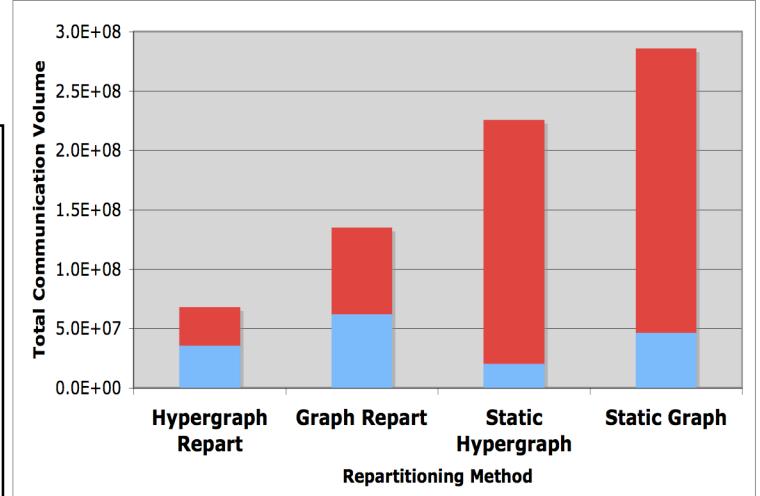
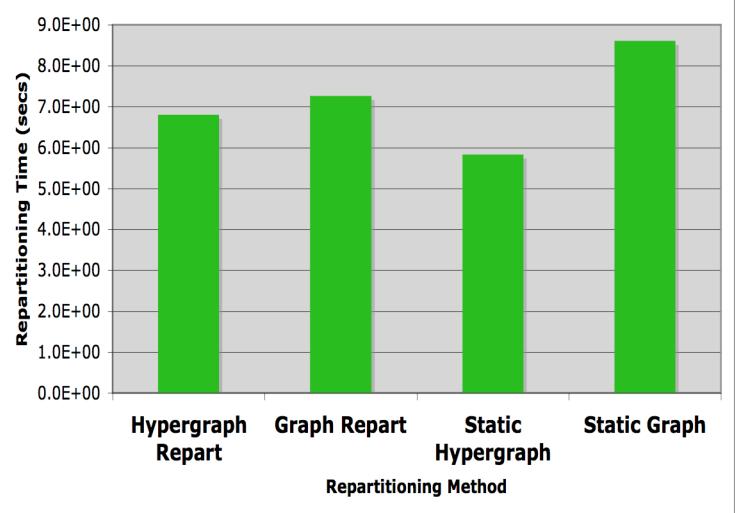
- Experiments with 64 parts on 64 processors.
- Dynamically adjust weights in data to simulate, say, adaptive mesh refinement.
- Repartition.
- Measure repartitioning time and total communication volume:

Data redistribution volume

+ Application communication volume

Total communication volume

# Repartitioning Results: Lower is Better

**SLAC 6.0M LCLS****Xyce 680K circuit****Repartitioning Time (secs)**



# Summary

---

- No one-size-fits-all solutions for partitioning.
- Different methods for different applications
  - Geometric vs. combinatorial/topological
  - Static vs. dynamic problem
- Zoltan toolkit has it all (almost...)
  - Provides collection of load-balance methods
  - Also provides other common parallel services
  - Frees the application developer to focus on his/her specialty area
  - Easy to test and compare different methods

## For More Information...

---

- **Zoltan Home Page**
  - <http://www.cs.sandia.gov/Zoltan>
  - User's and Developer's Guides
  - Download Zoltan software under GNU LGPL.
- **Email:**
  - [{egboman,kddevin}@sandia.gov](mailto:{egboman,kddevin}@sandia.gov)

---

# The End

---



# Example Hypergraph Callbacks

Slide 99



```
void ZOLTAN_HG_SIZE_CS_FN(void *data, int *num_lists, int *num_pins,  
    int *format, int *ierr);
```

Output from Application on Proc 0:

```
num_lists = 2  
num_pins = 6  
format = ZOLTAN_COMPRESSED_VERTEX  
        (owned non-zeros per vertex)  
ierr = ZOLTAN_OK
```

OR

Output from Application on Proc 0:

```
num_lists = 5  
num_pins = 6  
format = ZOLTAN_COMPRESSED_EDGE  
        (owned non-zeros per edge)  
ierr = ZOLTAN_OK
```

|   |  | Vertices |        |
|---|--|----------|--------|
|   |  | Proc 0   | Proc 1 |
|   |  | A        | B      |
| a |  | X        | X      |
| b |  | X        | X      |
| c |  |          | X X    |
| d |  | X        | X      |
| e |  | X        | X X    |
| f |  | X X      | X X    |

Hyperedges



# Example Hypergraph Callbacks

Slide 100



```
void ZOLTAN_HG_CS_FN(void *data, int num_gid_entries,  
int nvtxedge, int npins, int format,  
ZOLTAN_ID_PTR vtxedge_GID, int *vtxedge_ptr, ZOLTAN_ID_PTR pin_GID,  
int *ierr);
```

Proc 0 Input from Zoltan:

```
nvtxedge = 2 or 5  
npins = 6  
format = ZOLTAN_COMPRESSED_VERTEX or  
ZOLTAN_COMPRESSED_EDGE
```

Output from Application on Proc 0:

```
if (format = ZOLTAN_COMPRESSED_VERTEX)  
    vtxedge_GID = {A, B}  
    vtxedge_ptr = {0, 3}  
    pin_GID = {a, e, f, b, d, f}  
if (format = ZOLTAN_COMPRESSED_EDGE)  
    vtxedge_GID = {a, b, d, e, f}  
    vtxedge_ptr = {0, 1, 2, 3, 4}  
    pin_GID = {A, B, B, A, A, B}  
ierr = ZOLTAN_OK
```

|   |  | Vertices |        |
|---|--|----------|--------|
|   |  | Proc 0   | Proc 1 |
|   |  | A        | B      |
| a |  | X        | X      |
| b |  | X        | X      |
| c |  |          | X X    |
| d |  | X        | X      |
| e |  | X        | X X    |
| f |  | X X      | X X    |

Hyperedges